

# Decision Trees on MapReduce

CS246: Mining Massive Datasets  
Jure Leskovec, Stanford University  
<http://cs246.stanford.edu>




# Decision Tree Learning

- Give one attribute (e.g., lifespan), try to predict the value of new people's lifespans by means of some of the other available attribute
- **Input attributes:**
  - $d$  features/attributes:  $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots \mathbf{x}^{(d)}$
  - Each  $\mathbf{x}^{(j)}$  has **domain**  $O_j$ 
    - **Categorical:**  $O_j = \{\text{red, blue}\}$
    - **Numerical:**  $H_j = (0, 10)$
  - $Y$  is output variable with domain  $O_Y$ :
    - **Categorical:** Classification, **Numerical:** Regression
- **Data  $D$ :**
  - $n$  examples  $(\mathbf{x}_i, \mathbf{y}_i)$  where  $\mathbf{x}_i$  is a  $d$ -dim feature vector,  $\mathbf{y}_i \in O_Y$  is output variable
- **Task:**
  - Given an input data vector  $\mathbf{x}$  predict  $\mathbf{y}$

# Decision Trees

- A **Decision Tree** is a tree-structured plan of a set of attributes to test in order to predict the output




# Decision Trees (1)

## ■ Decision trees:

- Split the data at each internal node
- Each leaf node makes a prediction


## ■ Lecture today:

- Binary splits:  $X^{(j)} < v$
- Numerical attrrs.
- Regression



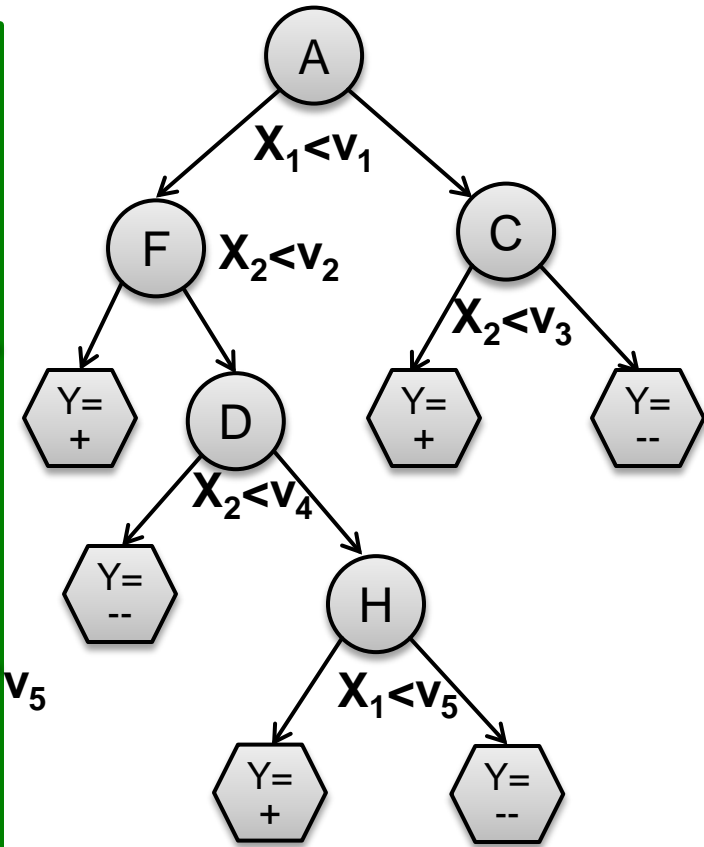

# How to make predictions?

- **Input:** Example  $\mathbf{x}_i$
- **Output:** Predicted  $y_i'$
- “Drop”  $\mathbf{x}_i$  down the tree until it hits a leaf node
- Predict the value stored in the leaf that  $\mathbf{x}_i$  hits



# Decision Trees Vs. SVM

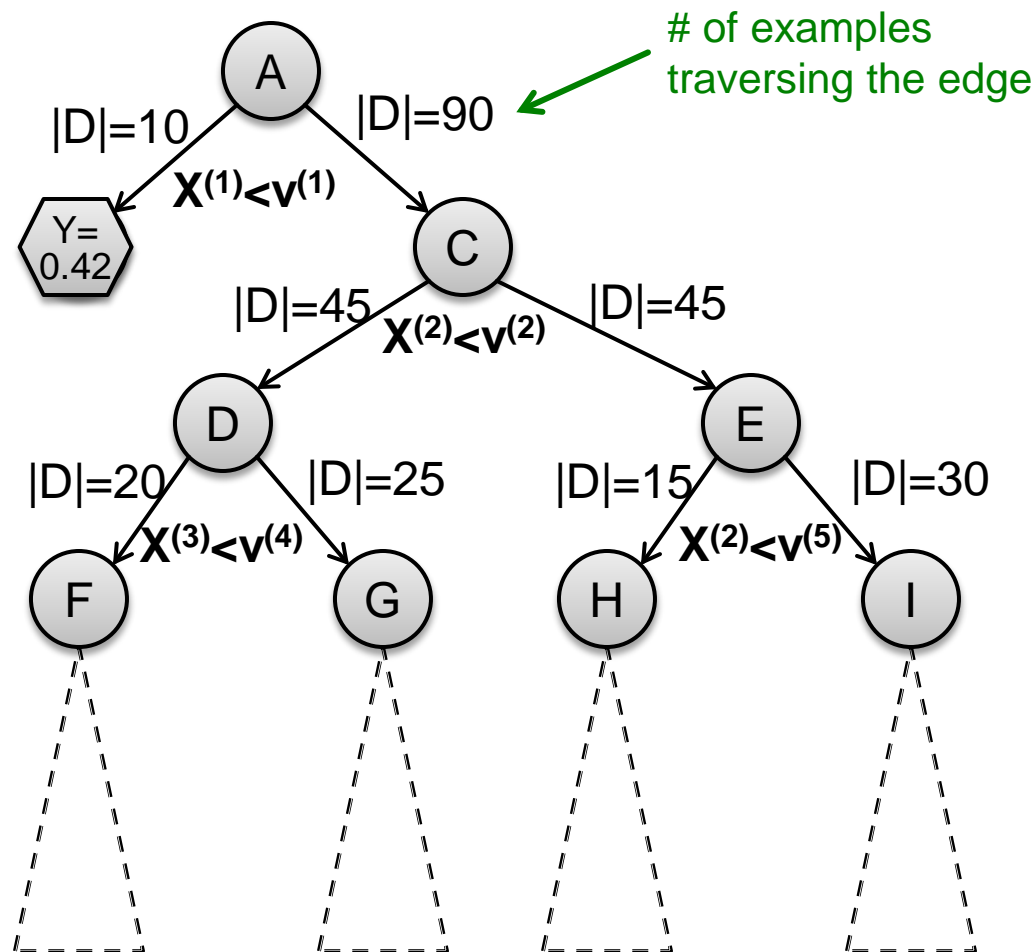
## ■ Alternative view:



**How to construct a tree?**

# How to construct a tree?


- Training dataset  $D^*$ ,  $|D^*|=100$  examples





# How to construct a tree?

- Imagine we are currently at some node  **$G$** 
  - Let  $D_G$  be the data that reaches  **$G$**
- **There is a decision we have to make: Do we continue building the tree?**
  - **If yes**, which variable and which value do we use for a **split**?
    - Continue building the tree recursively
  - **If not**, how do we make a prediction?
    - We need to build a “**predictor node**”



# 3 steps in constructing a tree

## Algorithm 1 BuildSubtree

Require: Node  $n$ , Data  $D \subseteq D^*$

1:  $(n \rightarrow \text{split}, D_L, D_R) = \text{FindBestSplit}(D)$  (1)

2: if  $\text{StoppingCriteria}(D_L)$  then (2)

3:    $n \rightarrow \text{left\_prediction} = \text{FindPrediction}(D_L)$  (3)

4: else

5:         **BuildSubtree** ( $n \rightarrow \text{left}, D_L$ )

6: if  $\text{StoppingCriteria}(D_R)$  then

7:    $n \rightarrow \text{right\_prediction} = \text{FindPrediction}(D_R)$

8: else

9:         **BuildSubtree** ( $n \rightarrow \text{right}, D_R$ )

- Requires at least a single pass over the data!

# How to construct a tree?


(1) How to split? Pick attribute & value that optimizes some criterion

■ Regression: Purity

- Find split  $(X^{(i)}, v)$  that creates  $D, D_L, D_R$ : parent, left, right child datasets and maximizes:

$$|D| \cdot \text{Var}(D) - (|D_L| \cdot \text{Var}(D_L) + |D_R| \cdot \text{Var}(D_R))$$

- $\text{Var}(D) = \frac{1}{n} \sum_{i \in D} (y_i - \bar{y})^2$  ... variance of  $y_i$  in  $D$




# How to construct a tree?

**(1) How to split?** Pick attribute & value that optimizes some criterion

- Classification:


**Information Gain**

- Measures how much a given attribute  $X$  tells us about the class  $Y$
- $IG(Y | X)$  : We must transmit  $Y$  over a binary link. How many bits on average would it save us if both ends of the line knew  $X$ ?



# Why Information Gain? Entropy

- **Entropy:** What's the smallest possible number of bits, on average, per symbol, needed to transmit a stream of symbols drawn from  $\mathbf{X}$ 's distribution?
- **The entropy of  $\mathbf{X}$ :**  $H(\mathbf{X}) = -\sum_{j=1}^m p_j \log p_j$ 
  - **“High Entropy”:**  $\mathbf{X}$  is from a uniform (boring) distribution
    - A histogram of the frequency distribution of values of  $\mathbf{X}$  is **flat**
  - **“Low Entropy”:**  $\mathbf{X}$  is from a varied (peaks/valleys) distrib.
    - A histogram of the frequency distribution of values of  $\mathbf{X}$  would have many lows and one or two highs



# Why Information Gain? Entropy

- Suppose I want to predict  $Y$  and I have input  $X$ 
  - $X$  = College Major
  - $Y$  = Likes “Gladiator”

X	Y
Math	Yes
History	No
CS	Yes
Math	No
Math	No
CS	Yes
Math	Yes
History	No

- From this data we estimate
  - $P(Y = \text{Yes}) = 0.5$
  - $P(X = \text{Math} \ \& \ Y = \text{No}) = 0.25$
  - $P(X = \text{Math}) = 0.5$
  - $P(Y = \text{Yes} \mid X = \text{History}) = 0$
- **Note:**
  - $H(Y) = -\frac{1}{2}\log_2(\frac{1}{2}) - \frac{1}{2}\log_2(\frac{1}{2}) = 1$
  - $H(X) = 1.5$

# Why Information Gain? Entropy

- Suppose I want to predict  $Y$  and I have input  $X$ 
  - $X$  = College Major
  - $Y$  = Likes “Gladiator”

$X$	$Y$
Math	Yes
History	No
CS	Yes
Math	No
Math	No
CS	Yes
Math	Yes
History	No

- Def: **Specific Conditional Entropy**
  - $H(Y \mid X=v)$  = The entropy of  $Y$  among only those records in which  $X$  has value  $v$
  - **Example:**
    - $H(Y \mid X = \text{Math}) = 1$
    - $H(Y \mid X = \text{History}) = 0$
    - $H(Y \mid X = \text{CS}) = 0$

# Why Information Gain?

- Suppose I want to predict  $Y$  and I have input  $X$ 
  - $X$  = College Major
  - $Y$  = Likes “Gladiator”

$X$	$Y$
Math	Yes
History	No
CS	Yes
Math	No
Math	No
CS	Yes
Math	Yes
History	No

- **Def: Conditional Entropy**

- $H(Y | X)$  = The average specific conditional entropy of  $Y$ 
  - = if you choose a record at random what will be the conditional entropy of  $Y$ , conditioned on that row's value of  $X$
  - = Expected number of bits to transmit  $Y$  if both sides will know the value of  $X$
- $= \sum_j P(X = v_j) H(Y | X = v_j)$



# Why Information Gain?

- Suppose I want to predict  $Y$  and I have input  $X$ 
  - $H(Y | X)$  = The average specific conditional entropy of  $Y$

X	Y
Math	Yes
History	No
CS	Yes
Math	No
Math	No
CS	Yes
Math	Yes
History	No

$$= \sum_j P(X = v_j) H(Y | X = v_j)$$

- Example:

$v_j$	$P(X=v_j)$	$H(Y X=v_j)$
Math	0.5	1
History	0.25	0
CS	0.25	0

- So:  $H(Y | X) = 0.5 * 1 + 0.25 * 0 + 0.25 * 0 = 0.5$

# Why Information Gain?

- Suppose I want to predict  $Y$  and I have input  $X$

- **Def: Information Gain**

- $IG(Y|X)$  = I must transmit  $Y$ . **How many bits on average would it save me if both ends of the line knew  $X$ ?**

$$IG(Y|X) = H(Y) - H(Y | X)$$

- **Example:**

- $H(Y) = 1$
- $H(Y|X) = 0.5$
- Thus  $IG(Y|X) = 1 - 0.5 = 0.5$

X	Y
Math	Yes
History	No
CS	Yes
Math	No
Math	No
CS	Yes
Math	Yes
History	No

# What is Information Gain used for?

- Suppose you are trying to predict whether someone is going live past 80 years
- From historical data you might find:
  - $IG(\text{LongLife} \mid \text{HairColor}) = 0.01$
  - $IG(\text{LongLife} \mid \text{Smoker}) = 0.3$
  - $IG(\text{LongLife} \mid \text{Gender}) = 0.25$
  - $IG(\text{LongLife} \mid \text{LastDigitOfSSN}) = 0.00001$
- IG tells us how much information about  $Y$  is contained in  $X$ 
  - So attribute  $X$  that has high  $IG(Y \mid X)$  is a good split!

# 3 steps in constructing a tree

## Algorithm 1 BuildSubtree

Require: Node  $n$ , Data  $D \subseteq D^*$

1:  $(n \rightarrow \text{split}, D_L, D_R) = \text{FindBestSplit}(D)$  (1)

2: if  $\text{StoppingCriteria}(D_L)$  then (2)

3:    $n \rightarrow \text{left\_prediction} = \text{FindPrediction}(D_L)$  (3)

4: else

5:         **BuildSubtree** ( $n \rightarrow \text{left}, D_L$ )

6: if  $\text{StoppingCriteria}(D_R)$  then

7:    $n \rightarrow \text{right\_prediction} = \text{FindPrediction}(D_R)$


8: else

9:         **BuildSubtree** ( $n \rightarrow \text{right}, D_R$ )

# When to stop?

## (2) When to stop?

- Many different heuristic options
- **Two ideas:**
  - **(1) When the leaf is “pure”**
    - The target variable does not vary too much:  $\text{Var}(y_i) < \varepsilon$
  - **(2) When # of examples in the leaf is too small**
    - For example,  $|D| \leq 100$



# How to predict?

## (3) How to predict?


### ■ Many options

#### ■ Regression:

- Predict average  $y_i$  of the examples in the leaf
- Build a linear regression model on the examples in the leaf

#### ■ Classification:

- Predict most common  $y_i$  of the examples in the leaf



# Building Decision Trees Using MapReduce

# Problem: Building a tree

- Given a large dataset with hundreds of attributes
- Build a decision tree!
- General considerations:
  - Tree is small (can keep it memory):
    - Shallow (~10 levels)
  - Dataset too large to keep in memory
  - Dataset too big to scan over on a single machine
  - MapReduce to the rescue!

---

**Algorithm 1 BuildSubTree**

---

Require: Node  $n$ , Data  $D \subseteq D^*$

```
1: ( $n \rightarrow \text{split}, D_L, D_R$ ) = FindBestSplit( $D$ )
2: if StoppingCriteria( $D_L$ ) then
3:    $n \rightarrow \text{left\_prediction}$  = FindPrediction( $D_L$ )
4: else
5:   BuildSubTree( $n \rightarrow \text{left}, D_L$ )
6: if StoppingCriteria( $D_R$ ) then
7:    $n \rightarrow \text{right\_prediction}$  = FindPrediction( $D_R$ )
8: else
9:   BuildSubTree( $n \rightarrow \text{right}, D_R$ )
```

---

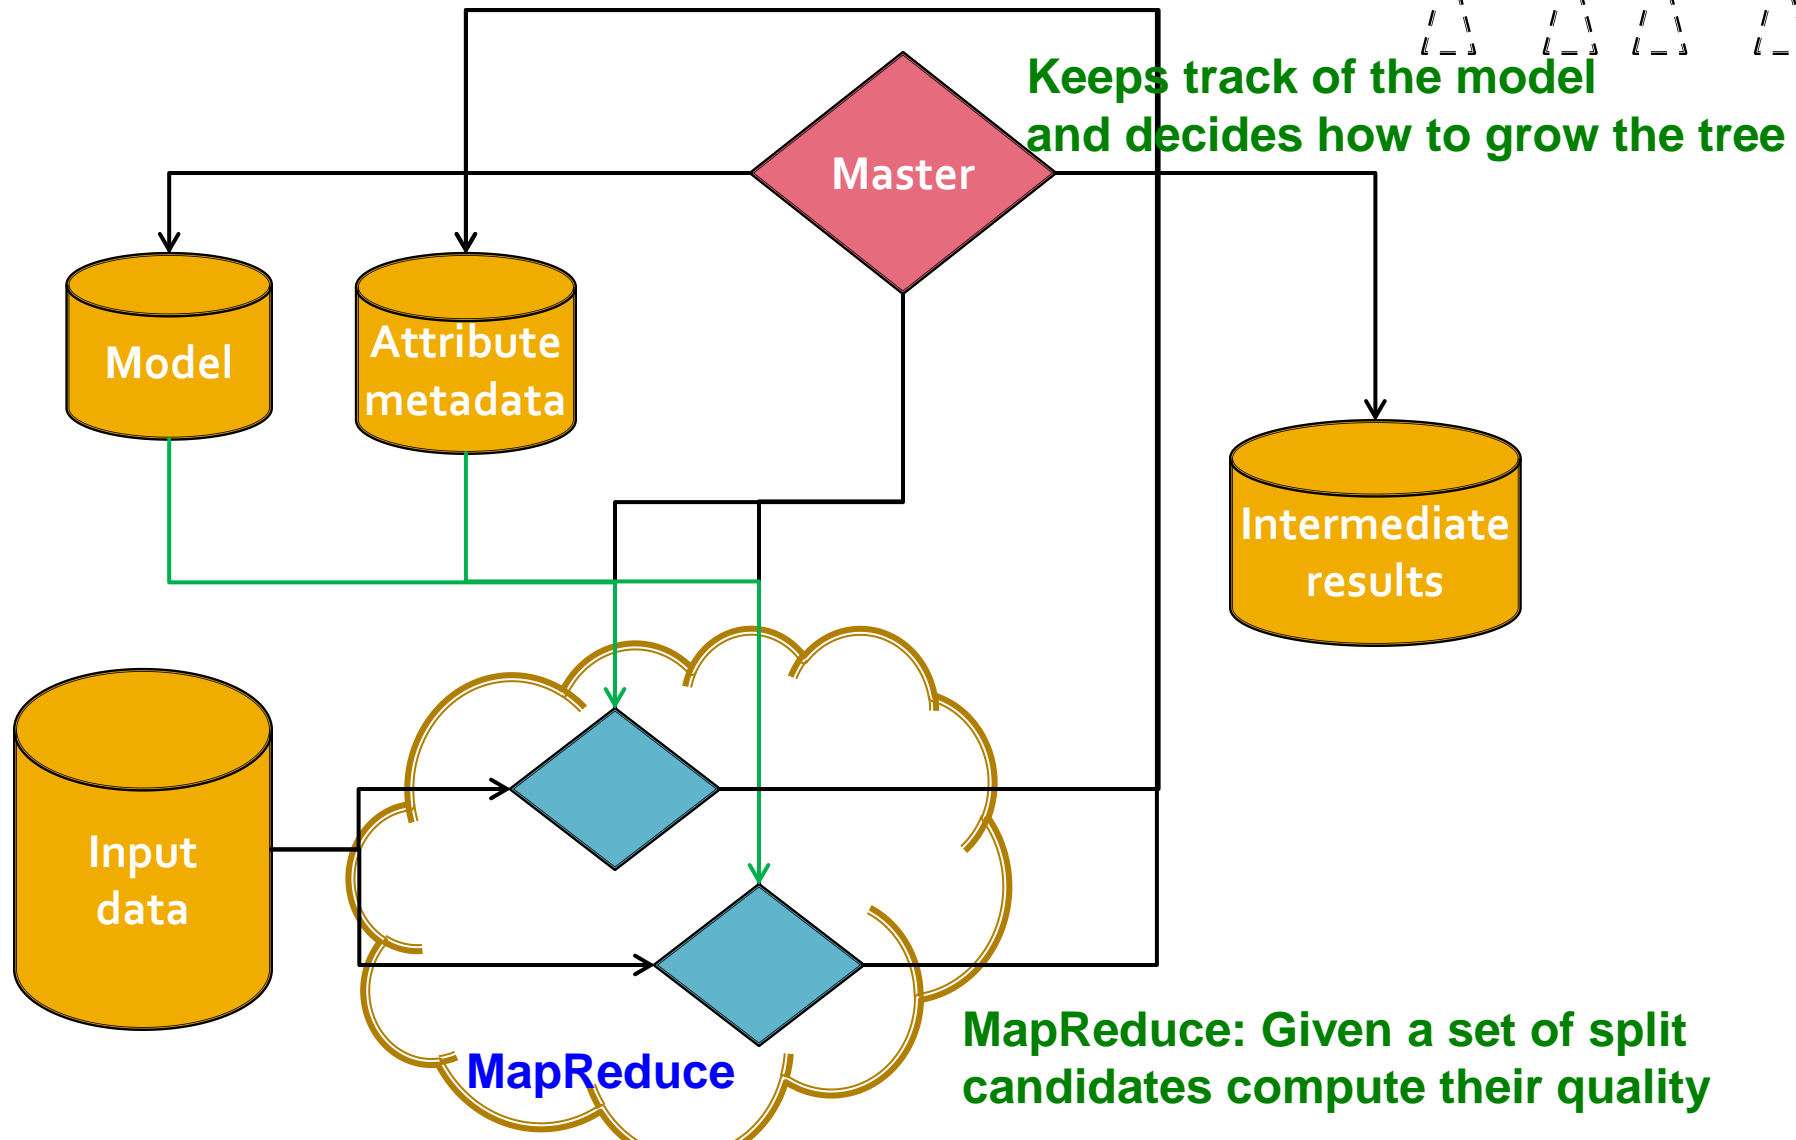



# Today's Lecture: PLANET

## Parallel Learner for Assembling Numerous Ensemble Trees [Panda et al., VLDB '09]

- A **sequence** of MapReduce jobs that builds a decision tree
- **Setting:**
  - Hundreds of **numerical** (discrete & continuous, but not categorical) attributes
  - Target variable is **numerical**: **Regression**
  - Splits are **binary**:  $X^{(i)} < v$
  - Decision tree is small enough for each Mapper to keep it in memory
  - Data too large to keep in memory



# PLANET Architecture



# PLANET: Building the Tree

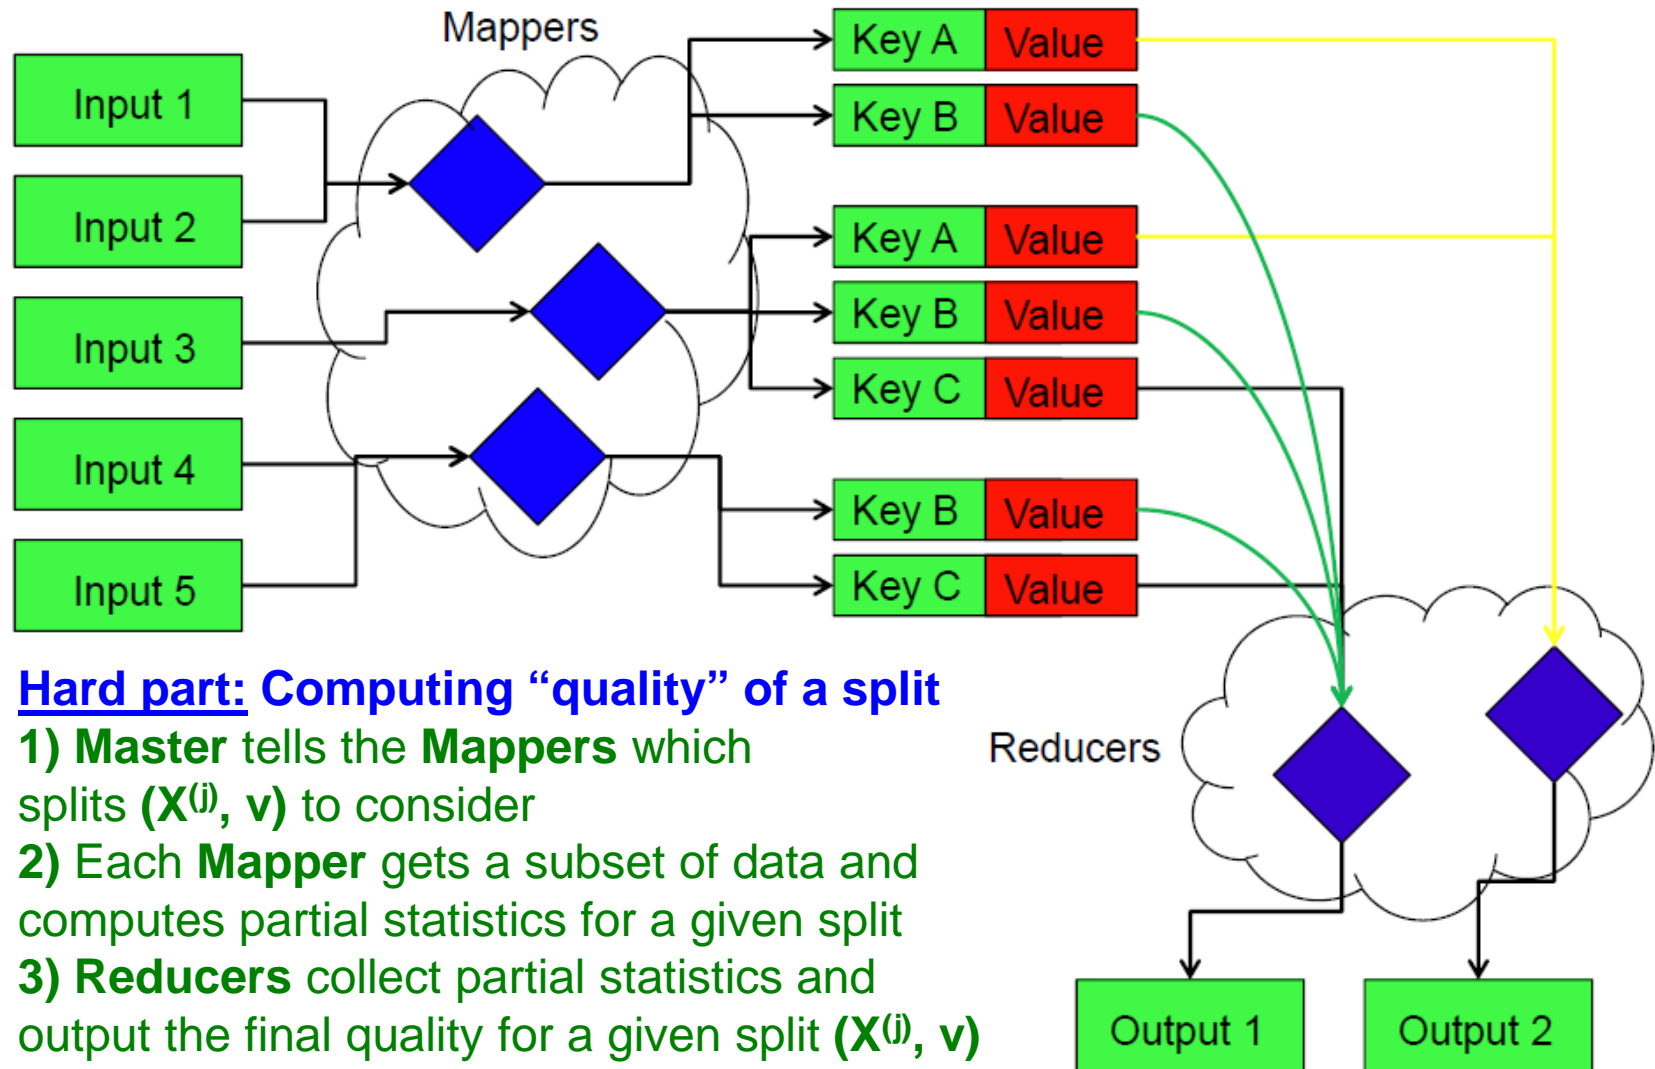
- **The tree will be built in levels**

- One level at a time:




- **1) Master decides which nodes/splits to consider, MapReduce computes quality of those splits**
- **2) Master then grows the tree for a level**
- **Goto 1)**

# Decision trees on MapReduce




# PLANET Overview




- We build the tree level by level
  - One MapReduce step builds **one level of the tree**
- **Mapper**
  - Considers a number of possible splits  $(X^{(i)}, v)$  on its subset of the data
  - For each split it stores **partial statistics**
  - Partial split-statistics is sent to **Reducers**
- **Reducer**
  - Collects all partial statistics and determines best split
- **Master** grows the tree for one level

# PLANET Overview



- **Mapper** loads the **model** and info about which **attribute splits** to consider
  - Each mapper sees a subset of the data  $D^*$
  - Mapper “drops” each datapoint to find the appropriate leaf node  $L$
  - For each leaf node  $L$  it keeps statistics about
    - (1) the data reaching  $L$
    - (2) the data in left/right subtree under split  $S$
- **Reducer** aggregates the statistics (1), (2) and determines the best split for each tree node

# PLANET: Components



- **Master**

- Monitors everything (runs multiple MapReduce jobs)

- **Three types of MapReduce jobs:**

- **(1) MapReduce Initialization (run once first)**

- For each attribute identify values to be considered for splits

- **(2) MapReduce FindBestSplit (run multiple times)**

- MapReduce job to find best split (when there is too much data to fit in memory)

- **(3) MapReduce InMemoryBuild (run once last)**

- Similar to **BuildSubTree** (but for small data)
- Grows an entire sub-tree once the data fits in memory

- **Model file**

- A file describing the state of the model

# PLANET: Components

(1) **Master Node**


(2) MapReduce Initialization (run once first)

(3) MapReduce FindBestSplit (run multiple times)

(4) MapReduce InMemoryBuild (run once last)



# PLANET: Master



- **Master controls the entire process**
- **Determines the state of the tree and grows it:**
  - (1) Decides if nodes should be split
  - (2) If there is little data entering a tree node, Master runs an InMemoryBuild MapReduce job to grow the entire subtree below that node
  - (3) For larger nodes, Master launches MapReduce FindBestSplit to evaluate candidates for best split
    - Master also collects results from **FindBestSplit** and chooses the best split for a node
  - (4) Updates the model

# PLANET: Components

- (1) Master Node
- (2) **MapReduce Initialization** (run once first)
- (3) MapReduce FindBestSplit (run multiple times)
- (4) MapReduce InMemoryBuild (run once last)

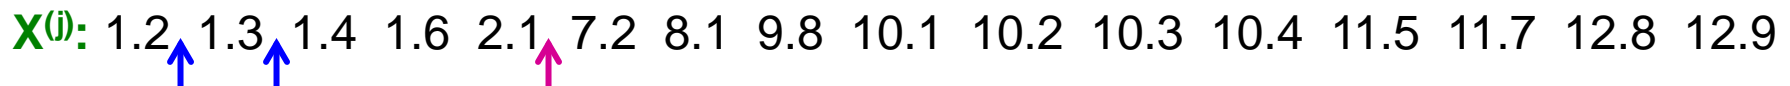
# Initialization: Attribute metadata

- **Initialization job:** Identifies all the attribute values which need to be considered for splits
  - Initialization process generates “**attribute metadata**” to be loaded in memory by other tasks
- **Main question:**  
**Which splits to even consider?**

# Initialization: Attribute metadata

- **Which splits to even consider?**
  - For small data we can sort the values along a particular feature and consider every possible split
  - **But data values may not be uniformly populated so many splits may not really make a difference**

$X^{(i)}$ : 1.2 1.3 1.4 1.6 2.1 7.2 8.1 9.8 10.1 10.2 10.3 10.4 11.5 11.7 12.8 12.9





The image shows a horizontal sequence of sorted data values: 1.2, 1.3, 1.4, 1.6, 2.1, 7.2, 8.1, 9.8, 10.1, 10.2, 10.3, 10.4, 11.5, 11.7, 12.8, 12.9. Three arrows point to specific values: a blue arrow points to 1.2, another blue arrow points to 1.3, and a pink arrow points to 2.1.

- **Idea:** Consider a limited number of splits such that splits “move” about the same amount of data

# Initialization: Attribute metadata


- **Splits for numerical attributes:**

- For attribute  $X^{(j)}$  we would like to consider every possible value  $v \in O_j$
- Compute an approx. equi-depth histogram on  $D^*$ 
  - **Idea:** Select buckets such that counts per bucket are equal



- Use boundary points of histogram as splits

# Side note: Computing Equi-Depth



- **Goal:** Equal number of elements per bucket ( $B$  buckets total)
- Construct by first **sorting** and then taking  **$B-1$**  equally-spaced splits

1 2 2 3 4 7 8 9 10 10 10 10 11 11 12 12 14 16 16 18 19 20 20 20

↑                    ↑                    ↑                    ↑                    ↑


- **Faster construction:**  
Sample & take equally-spaced splits in the sample
  - Nearly equal buckets

# PLANET: Components

- (1) Master Node
- (2) MapReduce Initialization (run once first)
- (3) MapReduce FindBestSplit (run multiple times)
- (4) MapReduce InMemoryBuild (run once last)

# FindBestSplit

- **Goal:** For a particular split node  $j$  find attribute  $X^{(j)}$  and value  $v$  that **maximizes Purity**:
  - $|D| \cdot Var(D) - (|D_L| \cdot Var(D_L) + |D_R| \cdot Var(D_R))$ 
    - $D$  ... training data  $(\mathbf{x}_i, y_i)$  reaching the node  $j$
    - $D_L$  ... training data  $\mathbf{x}_i$ , where  $\mathbf{x}_i^{(j)} < v$
    - $D_R$  ... training data  $\mathbf{x}_i$ , where  $\mathbf{x}_i^{(j)} \geq v$
    - $Var(D) = \frac{1}{n} \sum_{i \in D} y_i^2 - \left( \frac{1}{n} \sum_{i \in D} y_i \right)^2$






# FindBestSplit

- To compute Purity we need

- $Var(D) = \frac{1}{n} \sum_i y_i^2 - \left( \frac{1}{n} \sum_i y_i \right)^2$




- **Important observation:** Variance can be computed from **sufficient statistics**:

$$N, S = \sum y_i, Q = \sum y_i^2$$

- Each **Mapper** processes subset of data  $D_m$ , and computes  $N_m, S_m, Q_m$  for its own  $D_m$
  - **Reducer** combines the statistics and computes global variance and then Purity:

- $Var(D) = \frac{1}{\sum_m N_m} \sum_m Q_m - \left( \frac{1}{\sum_m N_m} \sum_m S_m \right)^2$


# FindBestSplit: Map



## ■ Mapper:

- Initialized by loading results of **Initialization task**
  - **Current model** (to find which node each datapoint  $x_i$  ends up)
  - **Attribute metadata** (all split points for each attribute)
- **For each data record run the Map algorithm:**
  - **For each node store statistics of the data entering the node and at the end emit (to all reducers):**
    - $\langle \text{NodeID}, \{ S=\sum y, Q=\sum y^2, N=\sum 1 \} \rangle$
  - **For each split store statistics and at the end emit:**
    - $\langle \text{SplitID}, \{ S, Q, N \} \rangle$ 
      - $\text{SplitID} = (\text{node } n, \text{attribute } X^{(j)}, \text{split value } v)$

# FindBestSplit: Reducer




## Reducer:

- (1) Load all the  $\langle \text{NodeID}, \underline{\text{List}} \{S_m, Q_m, N_m\} \rangle$  pairs and **aggregate** the per node statistics
- (2) For all the  $\langle \text{SplitID}, \underline{\text{List}} \{S_m, Q_m, N_m\} \rangle$  **aggregate** the statistics
  - $$\text{Var}(D) = \frac{1}{\sum_m N_m} \sum_m Q_m - \left( \frac{1}{\sum_m N_m} \sum_m S_m \right)^2$$
- **For each NodeID, output the best split found**

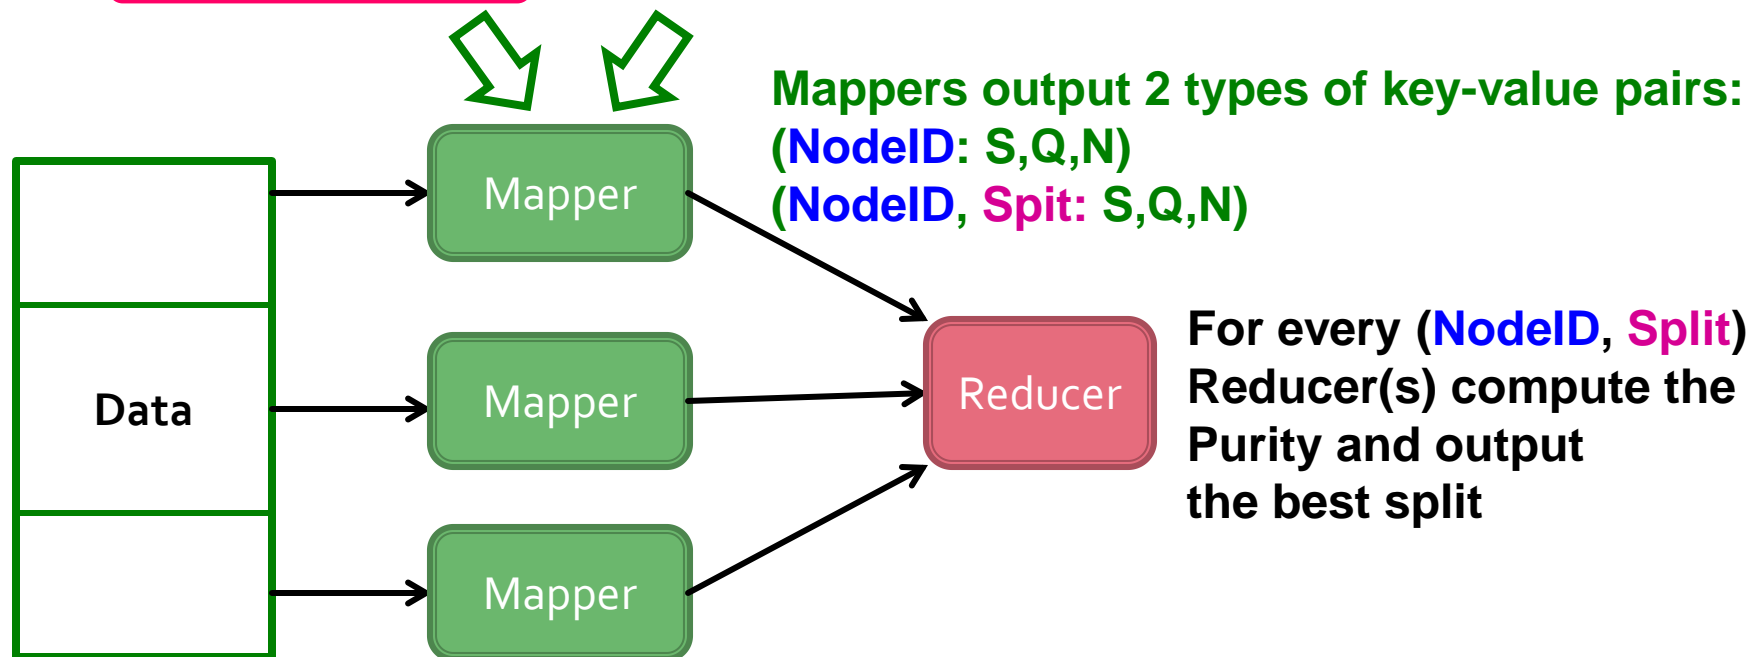
# Overall system architecture

- Master gives the mappers: (1) Tree  
(2) Set of nodes  
(3) Set of candidate splits




**Nodes:** F, G, H, I

**Split candidates:**  $(X^{(1)}, v^{(1)})$ ,  
 $(X^{(1)}, v^{(2)})$ ,  $(X^{(3)}, v^{(3)})$ ,  $(X^{(3)}, v^{(4)})$




# Overall system architecture


- **Example:** Need to split nodes **F, G, H, I**
- **Map and Reduce:**
  - **FindBestSplit::Map** (each mapper)
    - Load the current model **M**
    - Drop every example  $\mathbf{x}_i$  down the tree
    - If it hits **G** or **H**, update in-memory hash tables:
      - For each node:  $T_n: (\text{Node}) \rightarrow \{S, Q, N\}$
      - For each (Split, Node):  $T_{n,j,s}: (\text{Node}, \text{Attribute}, \text{SplitValue}) \rightarrow \{S, Q, N\}$
    - **Map::Finalize:** output the key-value pairs from above hashtables
  - **FindBestSplit::Reduce** (each reducer)
    - Collect:
      - $T1: \langle \text{Node}, \text{List}\{S, Q, N\} \rangle \rightarrow \langle \text{Node}, \{\Sigma S, \Sigma Q, \Sigma N\} \rangle$
      - $T2: \langle (\text{Node}, \text{Attr. Split}), \text{List}\{S, Q, N\} \rangle \rightarrow \langle (\text{Node}, \text{Attr. Split}), \{\Sigma S, \Sigma Q, \Sigma N\} \rangle$
    - Compute impurity for each node using **T1, T2**
    - Return **best split** to Master (which then decides on globally best split)



# Back to the Master



- Collects outputs from FindBestSplit reducers  
<Split.NodeID, Attribute, Value, Impurity>
- For each node decides the best split
  - If data in  $D_L/D_R$  is small enough, later run a MapReduce job **InMemoryBuild** on the node
  - Else run MapReduce **FindBestSplit** job for both nodes



# Decision Trees: Conclusion

# Decision Trees

- **Decision trees are the single most popular data mining tool:**
  - Easy to understand
  - Easy to implement
  - Easy to use
  - Computationally cheap
  - It's possible to get in trouble with overfitting
  - **They do classification as well as regression!**



# Learning Ensembles

- Learn multiple trees and combine their predictions
  - Gives better performance in practice
- Bagging:
  - Learns multiple trees over independent samples of the training data
    - For a dataset  $\mathbf{D}$  on  $n$  data points: Create dataset  $\mathbf{D}'$  of  $n$  points but sample from  $\mathbf{D}$  with replacement:
      - 33% points in  $\mathbf{D}'$  will be duplicates, 66% will be unique
  - Predictions from each tree are averaged to compute the final model prediction

# Bagged Decision Trees

- **How to create random samples of  $D^*$ ?**
  - Compute a hash of a training record's id and tree id
  - Use records that hash into a particular range to learn a tree
  - This way the same sample is used for all nodes in a tree
  - **Note:** This is sampling  $D^*$  without replacement (but samples of  $D^*$  should be created with replacement)

# SVM vs. DT

## ■ SVM

- **Classification**
  - Usually only 2 classes
- **Real valued features**  
(no categorical ones)
- **Tens/hundreds of thousands of features**
- **Very sparse features**
- **Simple decision boundary**
  - No issues with overfitting

## ■ Example applications

- Text classification
- Spam detection
- Computer vision

## ■ Decision trees

- **Classification & Regression**
  - Multiple (~10) classes
- **Real valued and categorical features**
- **Few (hundreds) of features**
- **Usually dense features**
- **Complicated decision boundaries**
  - Overfitting! Early stopping

## ■ Example applications

- User profile classification
- Landing page bounce prediction

# References

- B. Panda, J. S. Herbach, S. Basu, and R. J. Bayardo. ***PLANET: Massively parallel learning of tree ensembles with MapReduce***. In Proc. VLDB 2009.
- J. Ye, J.-H. Chow, J. Chen, Z. Zheng. ***Stochastic Gradient Boosted Distributed Decision Trees***. In Proc. CIKM 2009.