

- CS341 info session is on Thu 3/12 6pm in Gates415
- Final exam logistics
- Please fill out course evaluation forms (THANKS!!!)

Optimizing Submodular Functions

CS246: Mining Massive Datasets
Jure Leskovec, Stanford University
<http://cs246.stanford.edu>



Announcement: Final Exam Logistics

Final: At Stanford

- **Alternate final: Mon 3/16 7:00-10:00pm**
in Bishop Auditorium (Lathrop Library)
 - **Register at:** <http://goo.gl/forms/5505oC0Y94>
- **Final: Fri 3/20 12:15-3:15pm**
NVidia (Lastname starting with A-M)
Hewlett200 (Lastname starting with N-Z)
 - **See** <http://campus-map.stanford.edu>
 - **Practice finals + Gradiance quizzes will be on Piazza**
 - **Open book, open computer, no internet**
- **SCPD students can take the exam at Stanford!**

Final: SCPD Students

- **Exam protocol for SCPD students:**
 - On **Monday 3/16** your exam proctor will receive the PDF of the final exam from SCPD
 - **If you take the exam at Stanford:**
 - Ask the exam monitor to delete the SCPD email
 - **If you don't take the exam at Stanford:**
 - Arrange a **3h** slot with your exam monitor
 - You can take the exam **anytime** but return it in time
 - **Email exam PDF to cs246.mmds@gmail.com by Thursday 3/19 11:59pm Pacific time**

Announcement: CS341: Project in Mining Massive Datasets

- **Data mining research project on real data**
 - Groups of 3 students
 - **We provide interesting data, computing resources (Amazon EC2) and mentoring**
 - **You provide project ideas**
 - Class meets once a week + individual group mentoring

Information session:
Thursday 3/12 6:00pm in Gates 415
(there will be pizza!)

CS341: Schedule

- **Thu 3/12: Info session**
 - We will introduce datasets, problems, ideas
- **Students form groups and project proposals**
- **Mon 3/23: Project proposals are due**
- **We evaluate the proposals**
- **Mon 3/30: Admission results**
 - 10 to 15 groups/projects will be admitted
- **Mon 5/4, Wed 5/6: Midterm presentations**
- **Thu 6/11: Presentations, poster session**

More info: <http://cs341.stanford.edu>

Optimizing Submodular Functions

CS246: Mining Massive Datasets
Jure Leskovec, Stanford University
<http://cs246.stanford.edu>



Recommendations: Diversity

- Redundancy leads to a bad user experience

Obama Calls for Broad Action on Guns

**Obama unveils 23 executive actions,
calls for assault weapons ban**

**Obama seeks assault weapons ban,
background checks on all gun sales**

- Uncertainty around information need => don't put all eggs in one basket
- How do we optimize for diversity directly?

[illegible]

Hagel expects fight

10

[illegible]

New gun proposals

11

Encode Diversity as Coverage

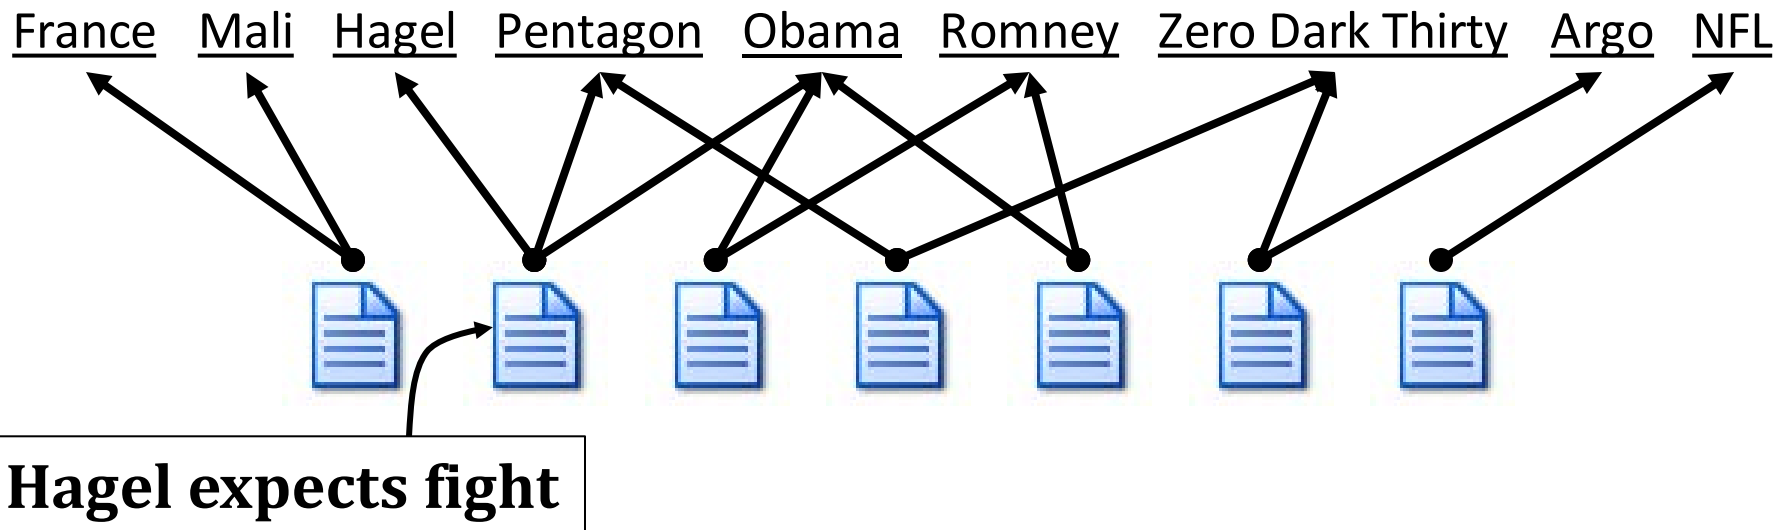
- **Idea:** Encode diversity as coverage problem
- **Example:** Word cloud of news for a single day
 - Want to select articles so that most words are “covered”



Diversity as Coverage

What is being covered?

- **Q: What is being covered?**
- **A: Concepts** (In our case: Named entities)



- **Q: Who is doing the covering?**
- **A: Documents**

Simple Abstract Model

- Suppose we are given a set of documents V
 - Each document d covers a set X_d of words/topics/named entities W

- For each set of documents A we define

$$F(A) = \left| \bigcup_{d \in A} X_d \right|$$


- Goal: We want to

$$\max_{|A| \leq k} F(A)$$

- Note: $F(A)$ is a set function: $F(A): \text{Sets} \rightarrow \mathbb{N}$

Maximum Coverage Problem

- Given universe of elements $W = \{w_1, \dots, w_n\}$ and sets $X_1, \dots, X_m \subseteq W$



- Goal: Find k sets X_i that cover the most of W
 - More precisely: Find k sets X_i whose size of the union is the largest
 - Bad news: A known NP-complete problem

Simple Greedy Heuristic

Simple Heuristic: Greedy Algorithm:

- Start with $A_0 = \{ \}$
- For $i = 1 \dots k$
 - Take set d that $\max F(A_{i-1} \cup \{d\})$
 - Let $A_i = A_{i-1} \cup \{d\}$


$$F(A) = \left| \bigcup_{d \in A} X_d \right|$$

■ Example:

- Eval. $F(\{d_1\}), \dots, F(\{d_m\})$, pick best (say d_1)
- Eval. $F(\{d_1\} \cup \{d_2\}), \dots, F(\{d_1\} \cup \{d_m\})$, pick best (say d_2)
- Eval. $F(\{d_1, d_2\} \cup \{d_3\}), \dots, F(\{d_1, d_2\} \cup \{d_m\})$, pick best
- And so on...


Simple Greedy Heuristic

- **Goal: Maximize the covered area**




Simple Greedy Heuristic

- **Goal: Maximize the covered area**




Simple Greedy Heuristic

- **Goal: Maximize the covered area**




Simple Greedy Heuristic

- **Goal: Maximize the covered area**




Simple Greedy Heuristic

- Goal: Maximize the covered area



When Greedy Heuristic Fails?



- **Goal:** Maximize the size of the covered area
- Greedy first picks A and then C
- But the optimal way would be to pick B and C

Approximation Guarantee

- Greedy produces a solution A
where: $F(A) \geq (1-1/e)*OPT$ ($F(A) \geq 0.63*OPT$)
[Nemhauser, Fisher, Wolsey '78]
- Claim holds for functions $F(\cdot)$ with 2 properties:
 - **F is monotone:** (adding more docs doesn't decrease coverage)
if $A \subseteq B$ then $F(A) \leq F(B)$ and $F(\{\})=0$
 - **F is submodular:**
adding an element to a set gives less improvement
than adding it to one of its subsets


Submodularity: Definition

Definition:

- Set function $F(\cdot)$ is called **submodular** if:

For all $A, B \subseteq W$:

$$F(A) + F(B) \geq F(A \cup B) + F(A \cap B)$$





Submodularity: Or equivalently

- **Diminishing returns** characterization

Equivalent definition:

- Set function $F(\cdot)$ is called **submodular** if:
For all $A \subseteq B$, $s \notin B$:

$$\underbrace{F(A \cup d) - F(A)}_{\text{Gain of adding } d \text{ to a small set}} \geq \underbrace{F(B \cup d) - F(B)}_{\text{Gain of adding } d \text{ to a large set}}$$



Example: Set Cover

- $F(\cdot)$ is **submodular**: $A \subseteq B$


$$\underbrace{F(A \cup d) - F(A)}_{\text{Gain of adding } X_d \text{ to a small set}} \geq \underbrace{F(B \cup d) - F(B)}_{\text{Gain of adding } X_d \text{ to a large set}}$$

Gain of adding X_d to a small set

Gain of adding X_d to a large set

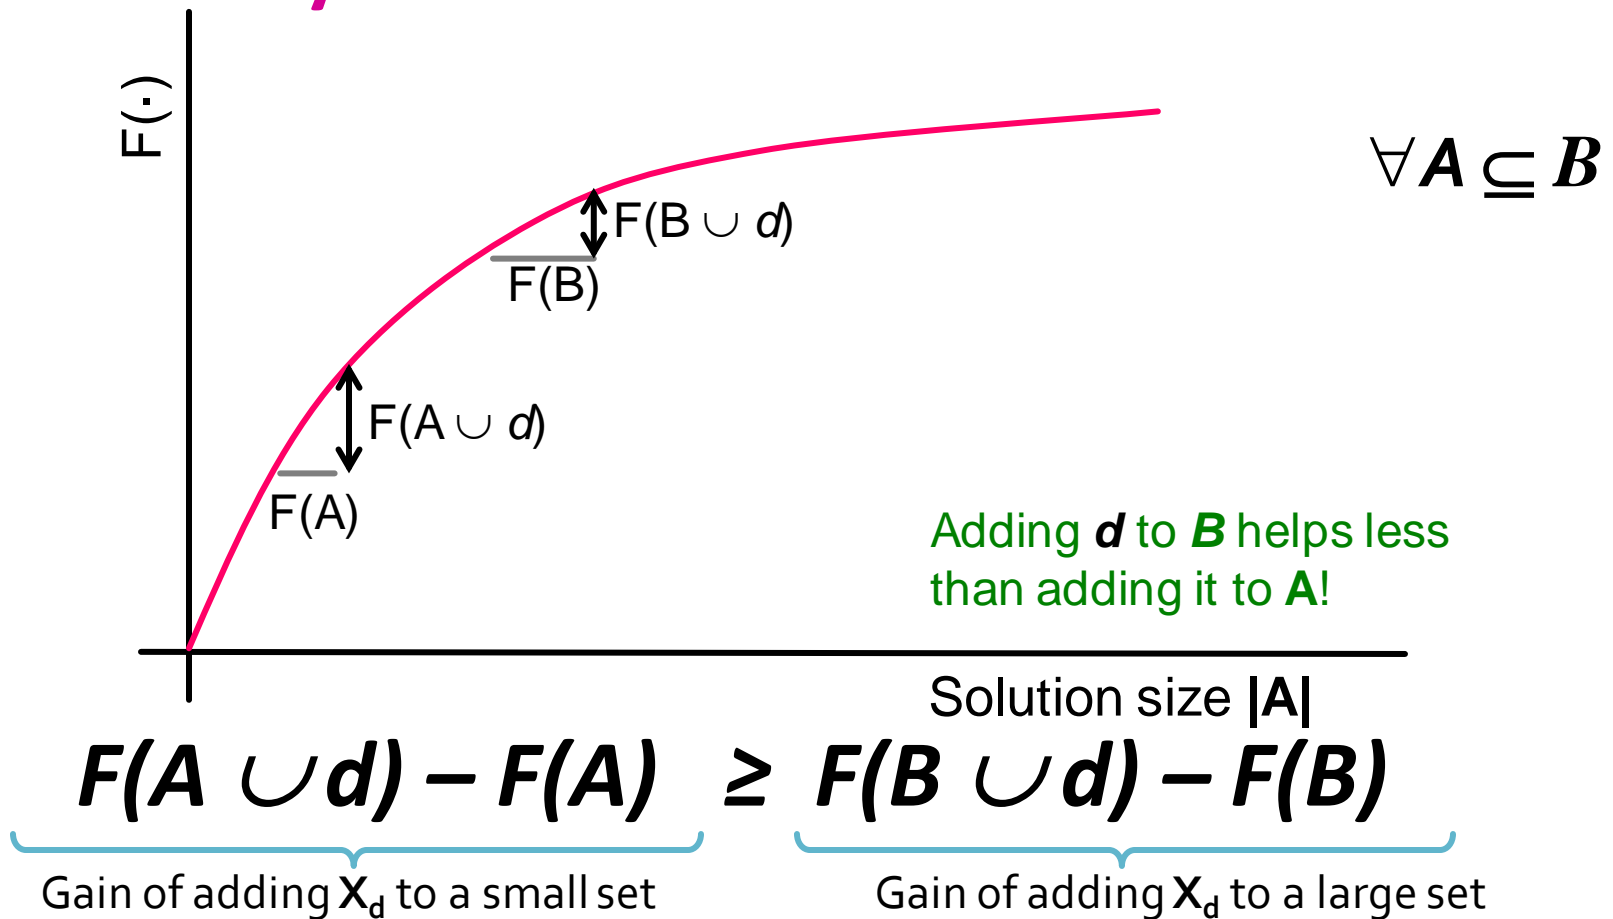
- **Natural example:**

- Sets X_1, \dots, X_m
- $F(A) = |\cup_{d \in A} X_d|$
(size of the covered area)
- Claim:
 $F(A)$ is submodular!



Submodularity– Diminishing returns

- Submodularity is discrete analogue of concavity



Submodularity & Concavity

- **Marginal gain:**

$$\Delta_F(d|A) = F(A \cup X_d) - F(A)$$

- **Submodular:**


$$A \subseteq B$$

$$F(A \cup d) - F(A) \geq F(B \cup d) - F(B)$$

- **Concavity:**

$$a \leq b$$

$$f(a + d) - f(a) \geq f(b + d) - f(b)$$

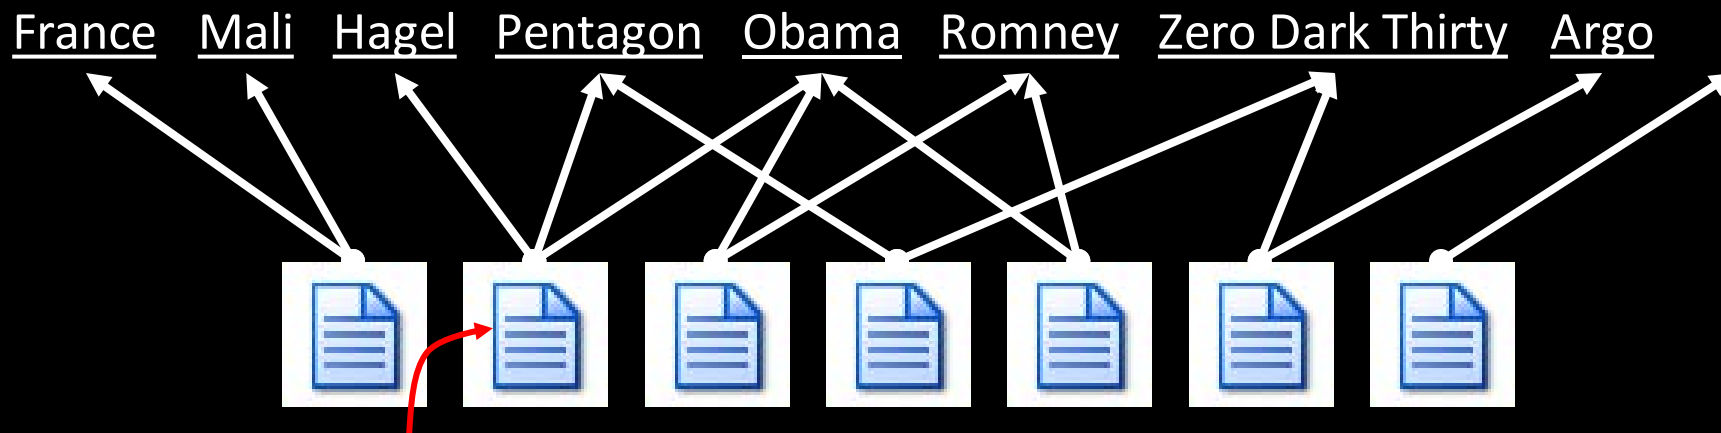


Submodularity: Useful Fact

- Let $F_1 \dots F_m$ be **submodular** and $\lambda_1 \dots \lambda_m > 0$
then $F(A) = \sum_i^m \lambda_i F_i(A)$ is **submodular**
 - **Submodularity is closed under non-negative linear combinations!**
- This is an extremely useful fact:
 - **Average of submodular functions is submodular:**
 $F(A) = \sum_i P(i) \cdot F_i(A)$
 - **Multicriterion optimization:** $F(A) = \sum_i \lambda_i F_i(A)$

Back to our problem

- **Q: What is being covered?**
- **A: Concepts** (In our case: Named entities)

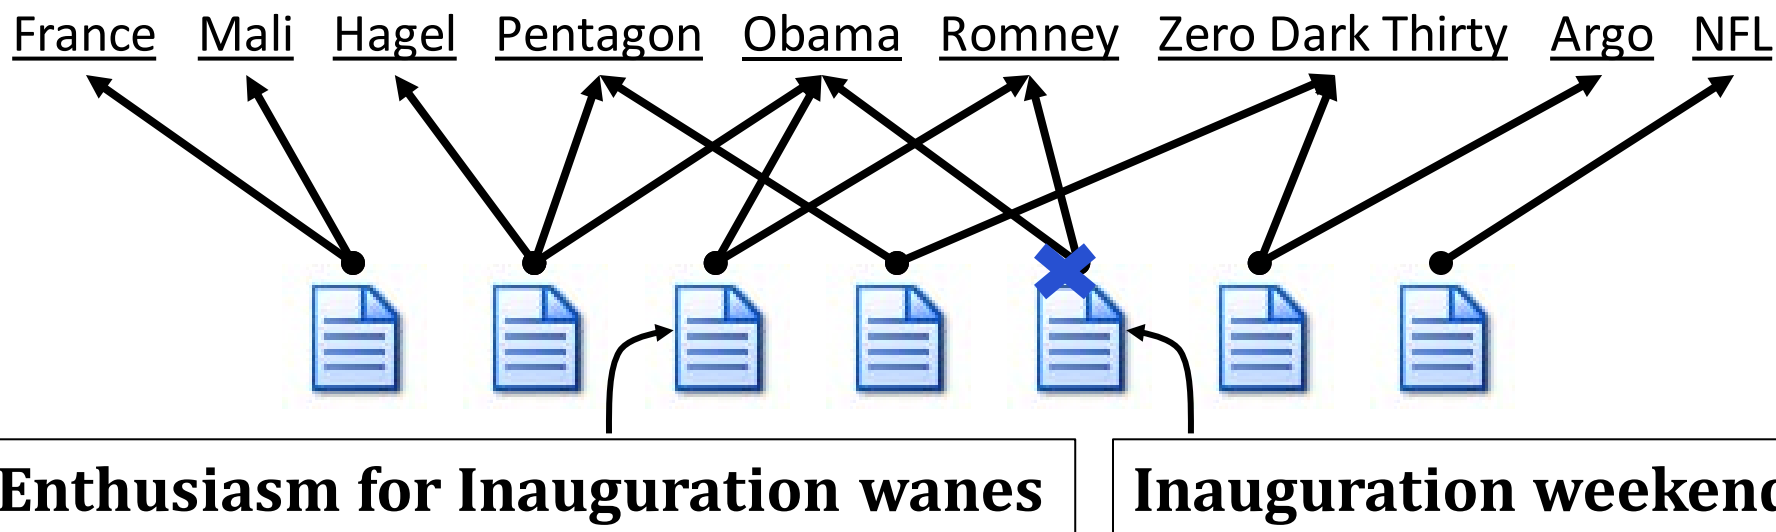


Hagel expects fight

- **Q: Who is doing the covering?**
- **A: Documents**

Back to our Concept Cover Problem

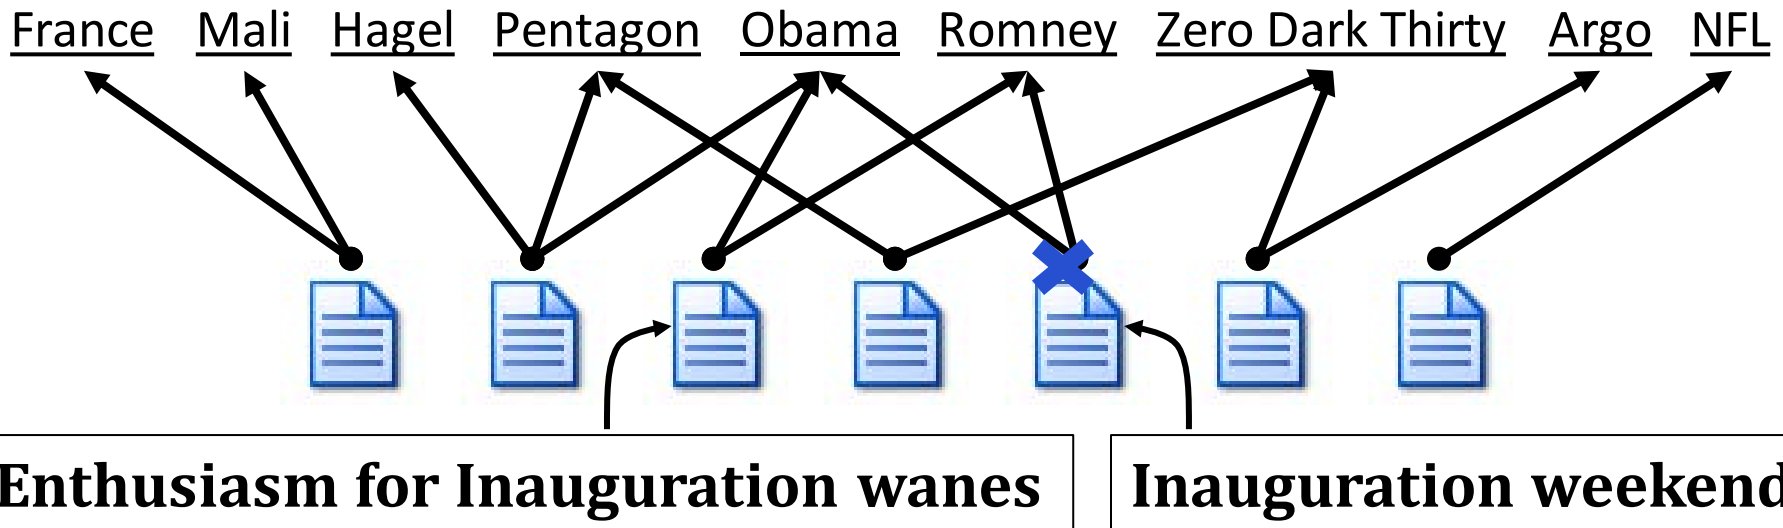
- **Objective:** pick k docs that cover most concepts



- $F(A)$: the number of concepts covered by A
 - *Elements...concepts, Sets ... concepts in docs*
 - $F(A)$ is submodular and monotone!
 - We can use **greedy** to optimize F

The Set Cover Problem

- **Objective:** pick k docs that cover most concepts



The good:

Penalizes redundancy

Submodular

The bad:

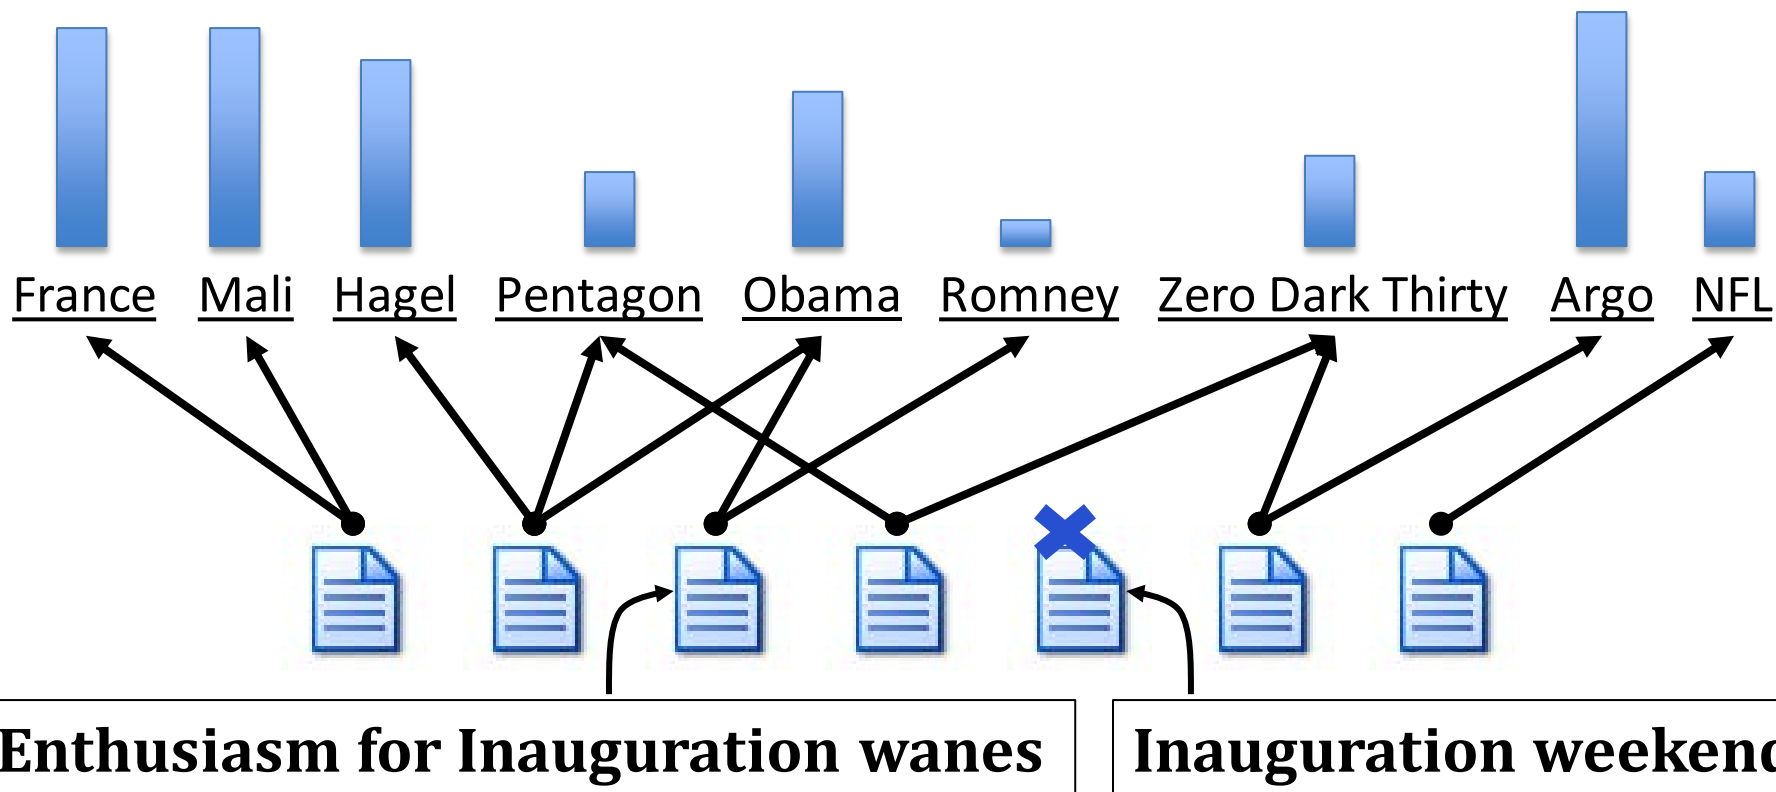
Concept importance?

All-or-nothing too harsh

Probabilistic Set Cover

Concept importance?

- **Objective:** pick k docs that cover most concepts

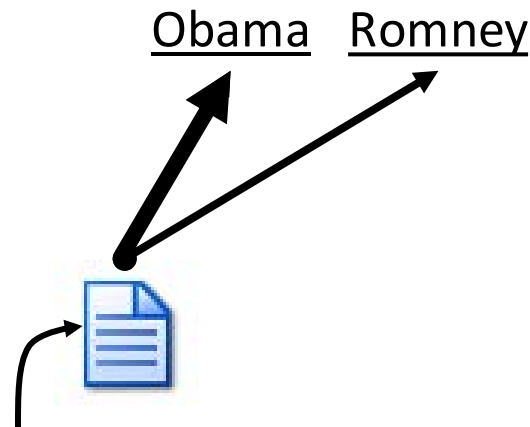


- Each concept c has importance weight w_c

All-or-nothing too harsh

- **Document coverage function**

$\text{cover}_d(c)$ = **probability** document **d** covers
concept **c**
[e.g., how strongly **d** covers **c**]



Enthusiasm for Inauguration wanes

Probabilistic Set Cover

- **Document coverage function:**

$\text{cover}_d(c) = \text{probability document } d \text{ covers concept } c$

- $\text{Cover}_d(c)$ can model how relevant is concept c for user u

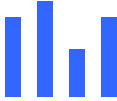
- **Set coverage function:**

$$\text{cover}_{\mathcal{A}}(c) = 1 - \prod_{d \in \mathcal{A}} (1 - \text{cover}_d(c))$$

- Prob. that at least one document in \mathbf{A} covers c

- **Objective:**

$$\max_{\mathcal{A}: |\mathcal{A}| \leq k} F(\mathcal{A}) = \sum_c w_c \text{cover}_{\mathcal{A}}(c)$$

concept weights 

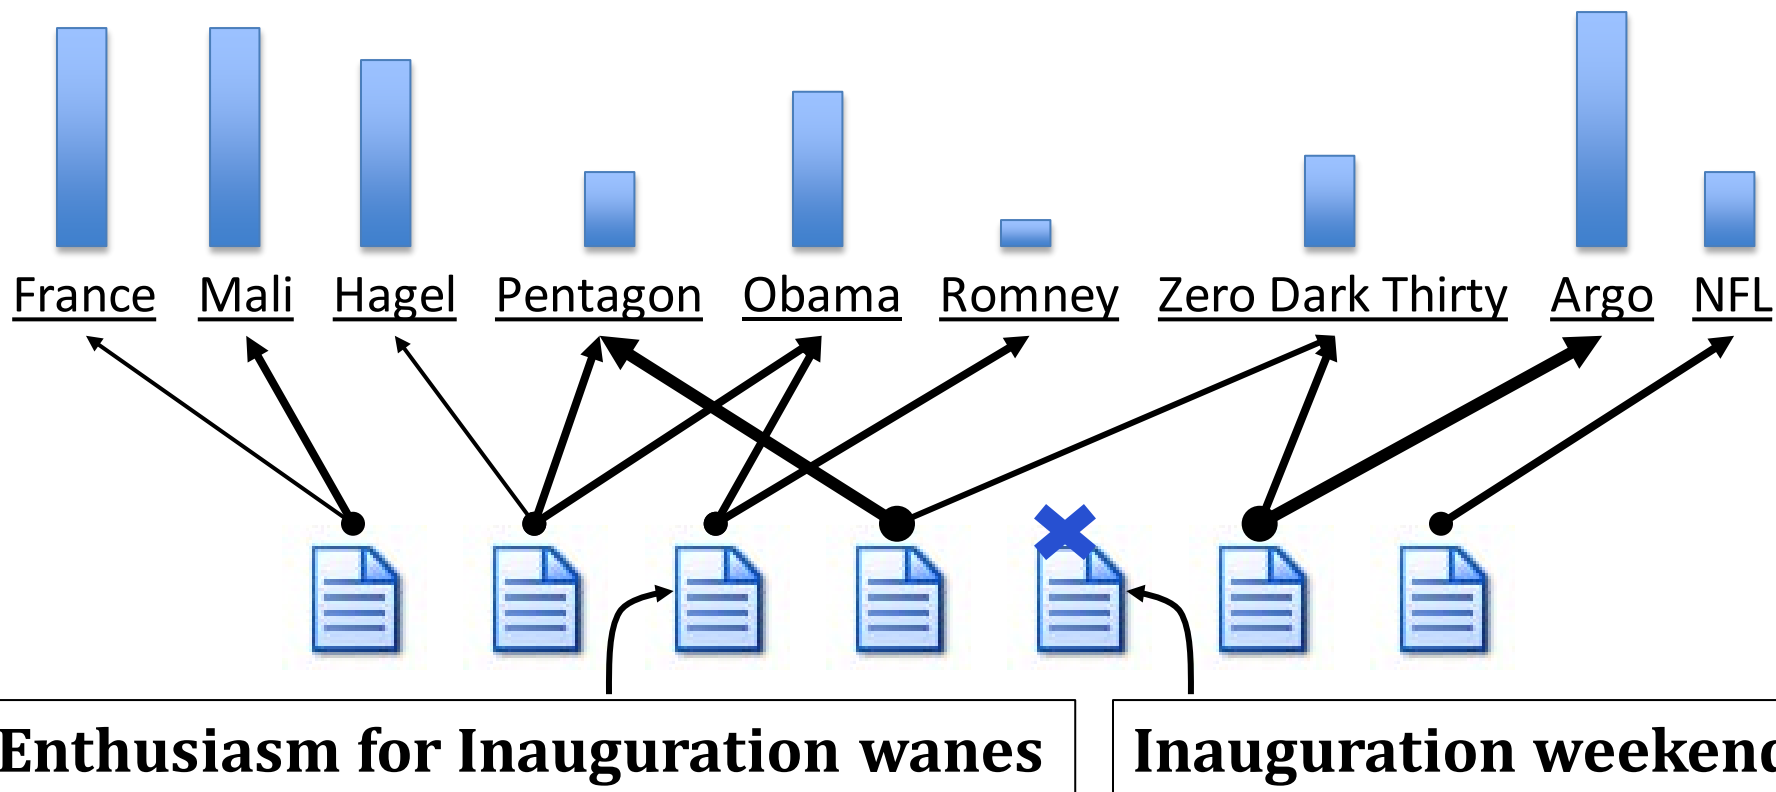
Optimizing $F(\mathcal{A})$

$$\max_{\mathcal{A}: |\mathcal{A}| \leq k} F(\mathcal{A}) = \sum_c w_c \text{cover}_{\mathcal{A}}(c)$$

- The objective function is also **submodular**
 - Intuitive **diminishing returns** property
 - Greedy algorithm leads to a $(1 - 1/e) \sim 63\%$ approximation, i.e., a **near-optimal** solution

Summary: Probabilistic Set Cover

- **Objective:** pick k docs that cover most concepts



- Each concept c has importance weight w_c
- Documents partially cover concepts: $\text{cover}_d(c)$




Lazy Optimization of Submodular Functions

Submodular Functions

Greedy

Marginal gain:

$$F(A \cup x) - F(A)$$



Add document with
highest marginal gain

- **Greedy algorithm is slow!**
 - At each iteration we need to re-evaluate marginal gains of **all remaining documents**
 - Runtime $O(|D| \cdot K)$ for selecting K documents out of the set of D of them

Speeding up Greedy

- **In round i :** So far we have $A_{i-1} = \{d_1, \dots, d_{i-1}\}$
 - Now we pick $d_i = \arg \max_{d \in V} F(A_{i-1} \cup \{d\}) - F(A_{i-1})$
 - Greedy algorithm maximizes the “marginal benefit”

$$\Delta_i(d) = F(A_{i-1} \cup \{d\}) - F(A_{i-1})$$

- **By submodularity property:**

$$F(A_i \cup \{d\}) - F(A_i) \geq F(A_j \cup \{d\}) - F(A_j) \text{ for } i < j$$


- **Observation: By submodularity:**

For every $d \in D$

$$\Delta_i(d) \geq \Delta_j(d) \text{ for } i < j \text{ since } A_i \subseteq A_j$$

$$\Delta_i(d) \geq \Delta_j(d)$$

- **Marginal benefits $\Delta_i(d)$ only shrink!**
(as i grows)



Selecting document d in step i covers more words than selecting d at step j ($j > i$)

Lazy Greedy


■ Idea:

- Use Δ_i as upper-bound on Δ_j ($j > i$)

■ Lazy Greedy:

- Keep an ordered list of marginal benefits Δ_i from previous iteration
- Re-evaluate Δ_i **only** for top node
- Re-sort and prune

(Upper bound on)
Marginal gain Δ_1



$$F(A \cup \{d\}) - F(A) \geq F(B \cup \{d\}) - F(B) \quad A \subseteq B$$

Lazy Greedy


■ Idea:

- Use Δ_i as upper-bound on Δ_j ($j > i$)

■ Lazy Greedy:

- Keep an ordered list of marginal benefits Δ_i from previous iteration
- Re-evaluate Δ_i **only** for top node
- Re-sort and prune

Upper bound on
Marginal gain Δ_2



$$F(A \cup \{d\}) - F(A) \geq F(B \cup \{d\}) - F(B) \quad A \subseteq B$$

Lazy Greedy


■ Idea:

- Use Δ_i as upper-bound on Δ_j ($j > i$)

■ Lazy Greedy:

- Keep an ordered list of marginal benefits Δ_i from previous iteration
- Re-evaluate Δ_i **only** for top node
- Re-sort and prune

Upper bound on
Marginal gain Δ_2




$$F(A \cup \{d\}) - F(A) \geq F(B \cup \{d\}) - F(B) \quad A \subseteq B$$

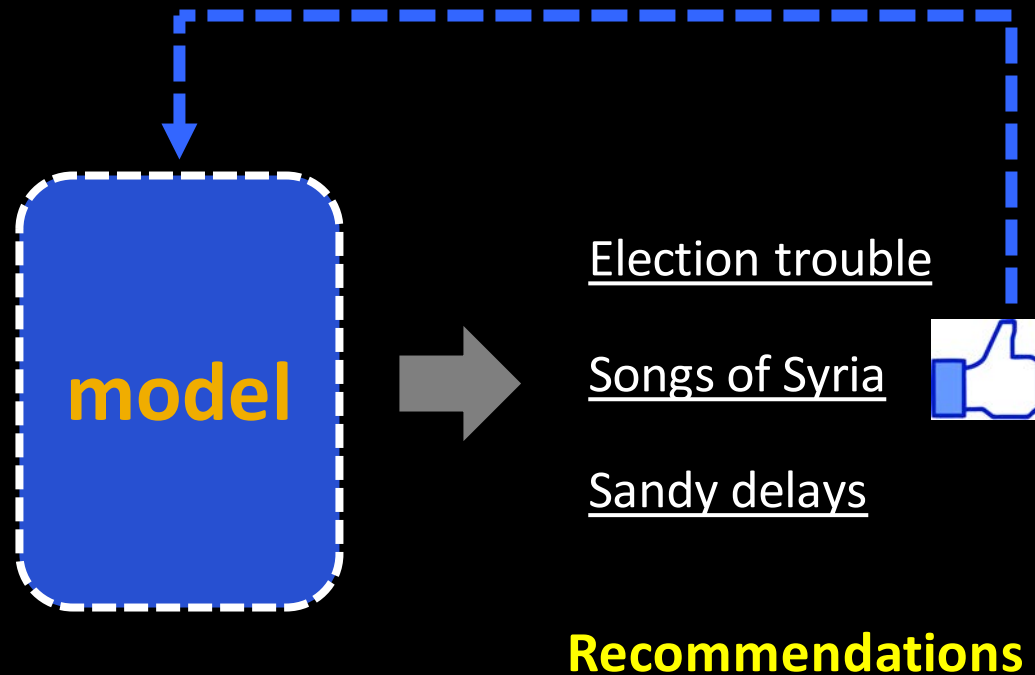
Summary so far

■ Summary so far:

- Diversity can be formulated as a set cover
- Set cover is submodular optimization problem
- Can be (approximately) solved using greedy algorithm
- Lazy-greedy gives significant speedup

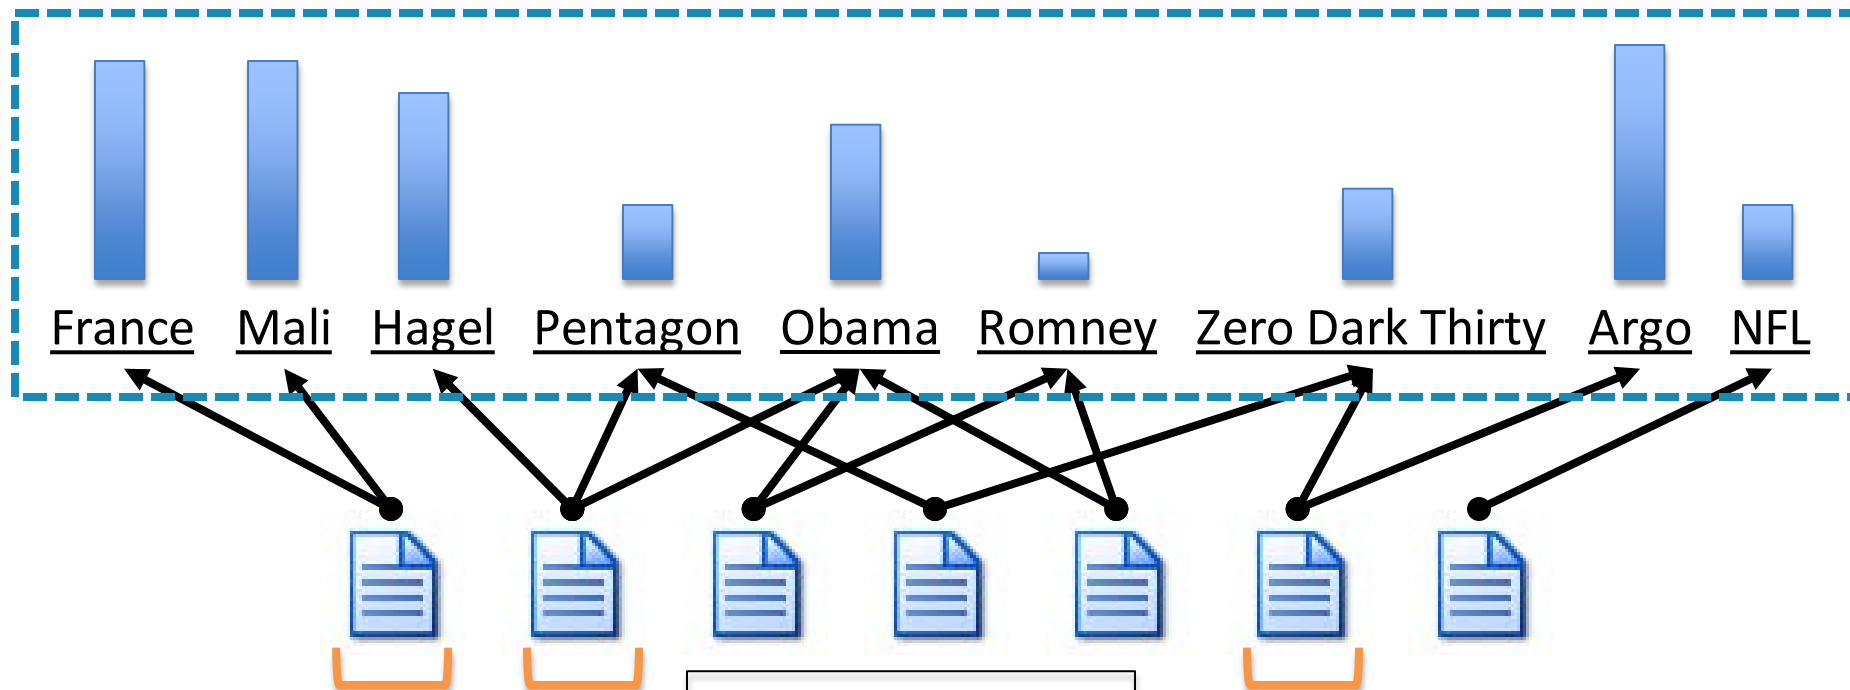


But what about **personalization?**



Concept Coverage

We assumed same concept weighting for all users



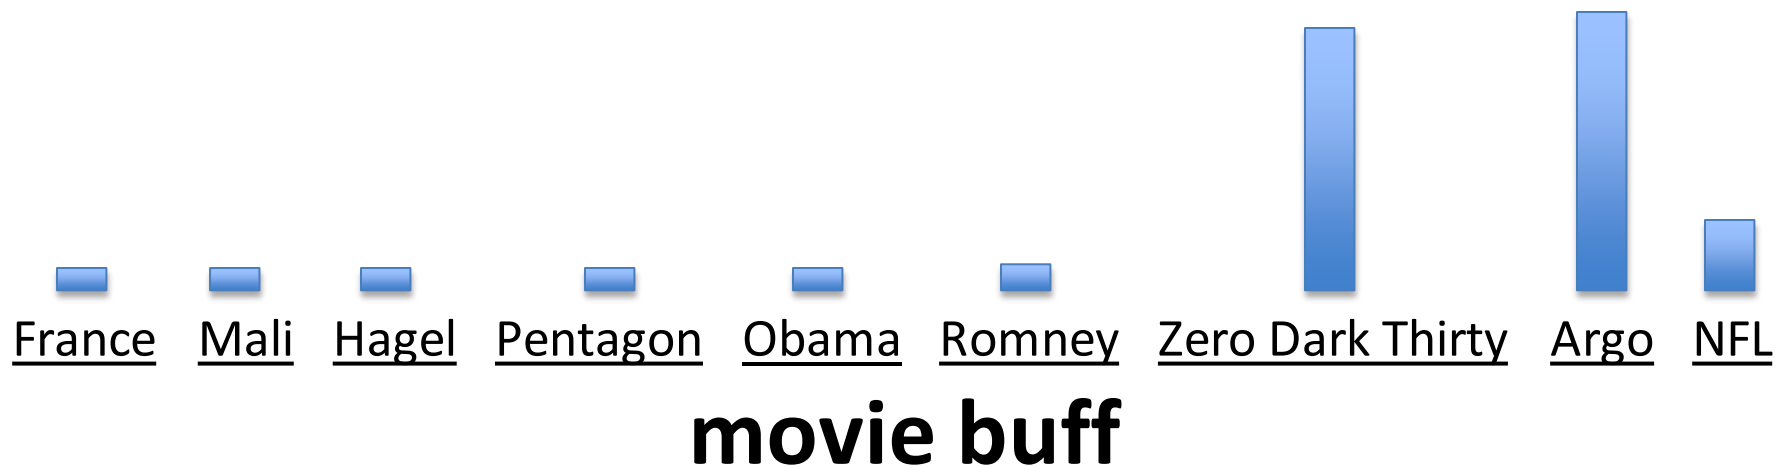

France intervenes

Chuck for Defense

Argo wins big

Personal Concept Weights

- Each user has **different** preferences over concepts



Personal concept weights

- Assume each user u has **different** preference vector $\mathbf{w}_c^{(u)}$ over concepts c

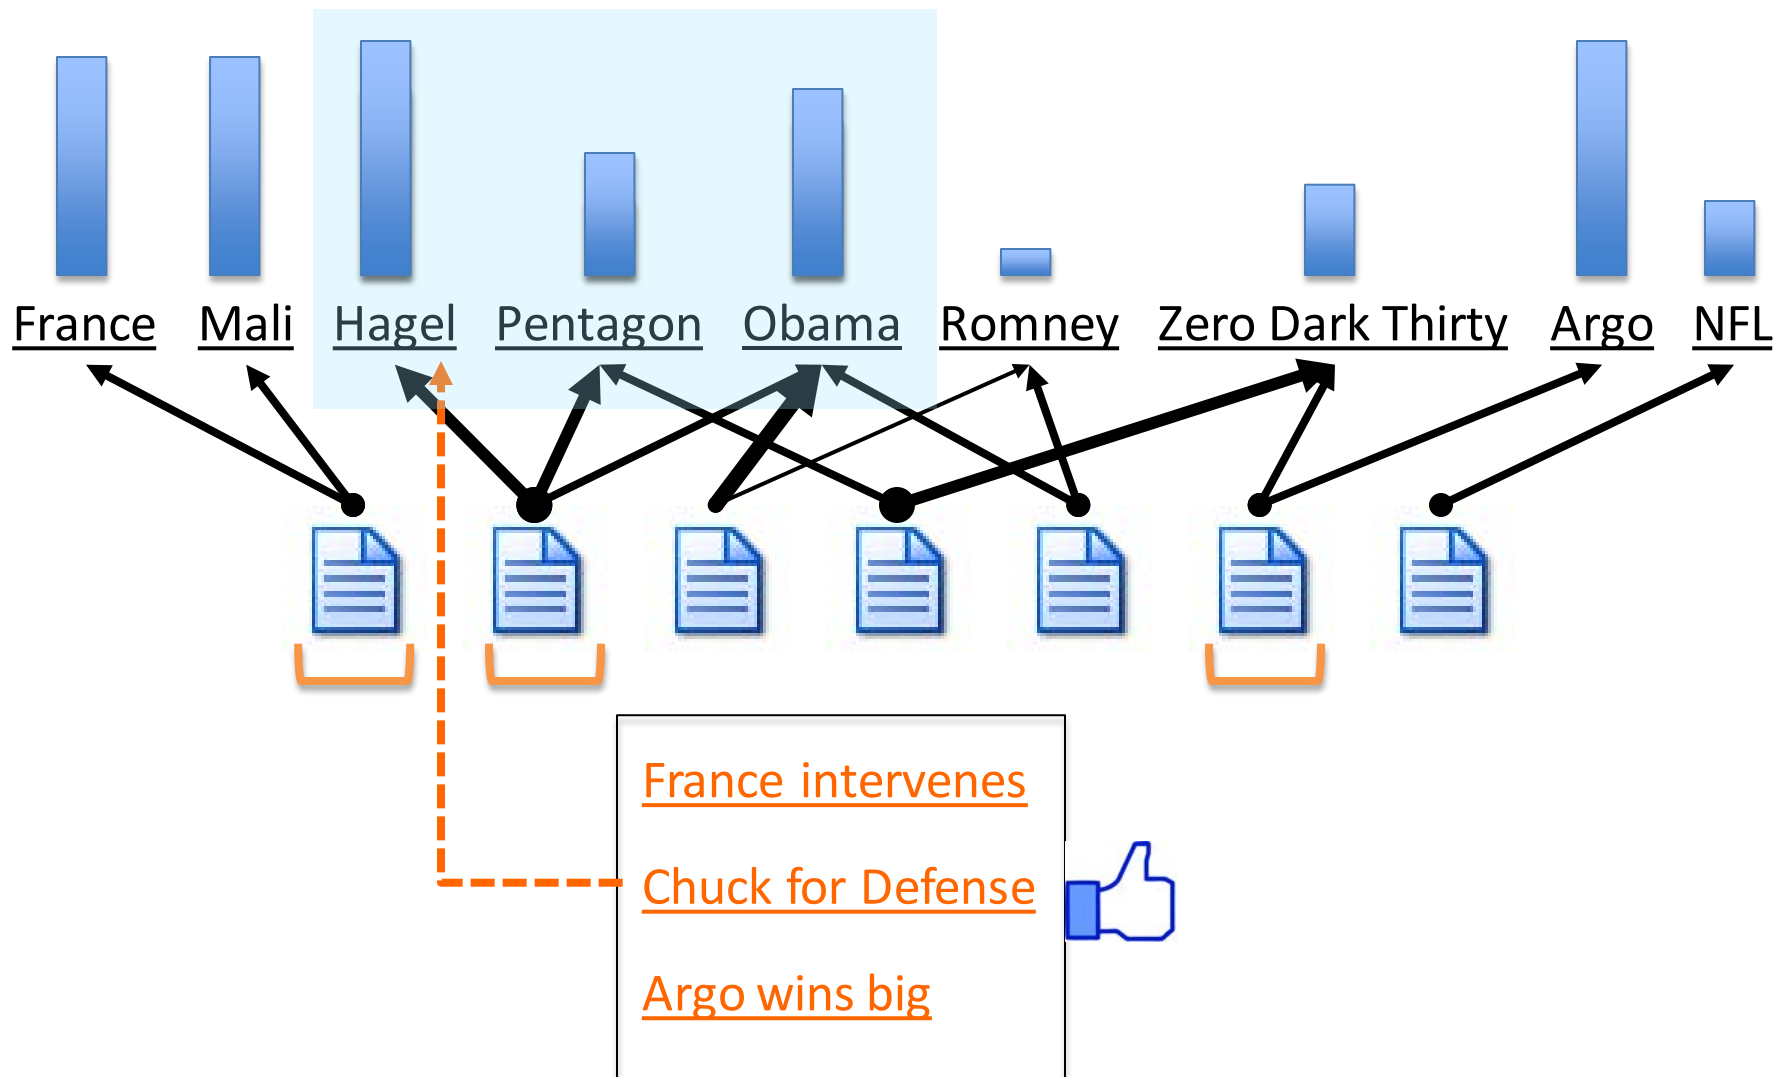
$$\max_{\mathcal{A}: |\mathcal{A}| \leq k} F(\mathcal{A}) = \sum_c w_c \text{cover}_{\mathcal{A}}(c)$$



$$\max_{\mathcal{A}: |\mathcal{A}| \leq k} F(\mathcal{A}) = \sum_c w_c^{(u)} \text{cover}_{\mathcal{A}}(c)$$

- Goal:** Learn personal concept weights from user feedback

Interactive Concept Coverage



Multiplicative Weights (MW)

- **Multiplicative Weights algorithm**
 - Assume each concept c has weight w_c
 - We recommend document d and receive feedback, say $r = +1$ or -1
 - **Update the weights:**
 - If $c \in X_d$ then $w_c = \beta^r w_c$
 - If $c \notin X_d$ then $w_c = \beta^{-r} w_c$
 - If concept c appears in X_d and we received positive feedback $r=+1$ then we increase the weight w_c by multiplying it by β ($\beta > 1$) otherwise we decrease the weight (divide by β)
 - **Normalize weights so that $\sum_c w_c = 1$**

Summary of the Algorithm

■ Steps of the algorithm:

1. Identify **items** to recommend from
2. Identify **concepts** [what makes items redundant?]
3. **Weigh** concepts by general importance
4. Define **item-concept coverage function**
5. **Select** items using probabilistic set cover
6. Obtain **feedback**, **update** weights

Summary: Submodularity

| | Maximization | Minimization |
|---------------|--|--|
| Unconstrained | NP-hard, but well-approximable (if nonnegative) | Polynomial time! Generally inefficient (n^6), but can exploit special cases (cuts; symmetry; decomposable; ...) |
| Constrained | NP-hard but well-approximable „Greedy-(like)“ for cardinality, matroid constraints; Non-greedy for more complex (e.g., connectivity) constraints | NP-hard; hard to approximate, still useful algorithms |