

Mining Massive Datasets: Review

CS246: Mining Massive Datasets
Jure Leskovec, Stanford University
<http://cs246.stanford.edu>

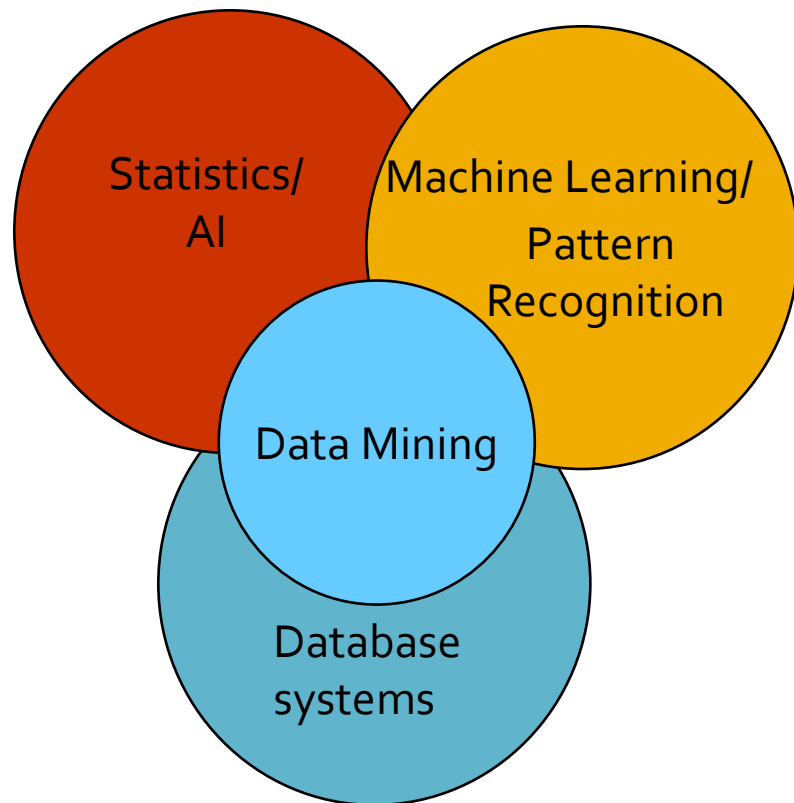


Data Mining

- **Models and tools for discovering patterns and answering queries that are:**
 - **Valid:** Hold on new data with some certainty
 - **Useful:** Should be possible to act on the item
 - **Unexpected:** Non-obvious to the system
 - **Understandable:** Humans should be able to interpret the pattern

Mining Massive Datasets

- Overlaps with machine learning, statistics, artificial intelligence, databases, but more stress on
 - **Scalability** of number of features and instances
 - **Algorithms** and **architectures**
 - Automation for handling **large data**



What We Have Covered

- Apriori
- MapReduce
- Association rules
- Frequent itemsets
- PCY
- Recommender systems
- PageRank
- TrustRank
- HITS
- SVM
- Decision Trees
- Perceptron
- Web Advertising
- DGIM
- Bandits
- BFR
- Regret
- LSH
- MinHash
- SVD
- Clustering
- Matrix factorization
- CUR
- Bloom filters
- Flajolet-Martin
- CURE
- Submodularity
- SGD
- Collaborative Filtering
- SimRank
- Random hyperplanes
- Trawling
- AND-OR constructions
- k-means

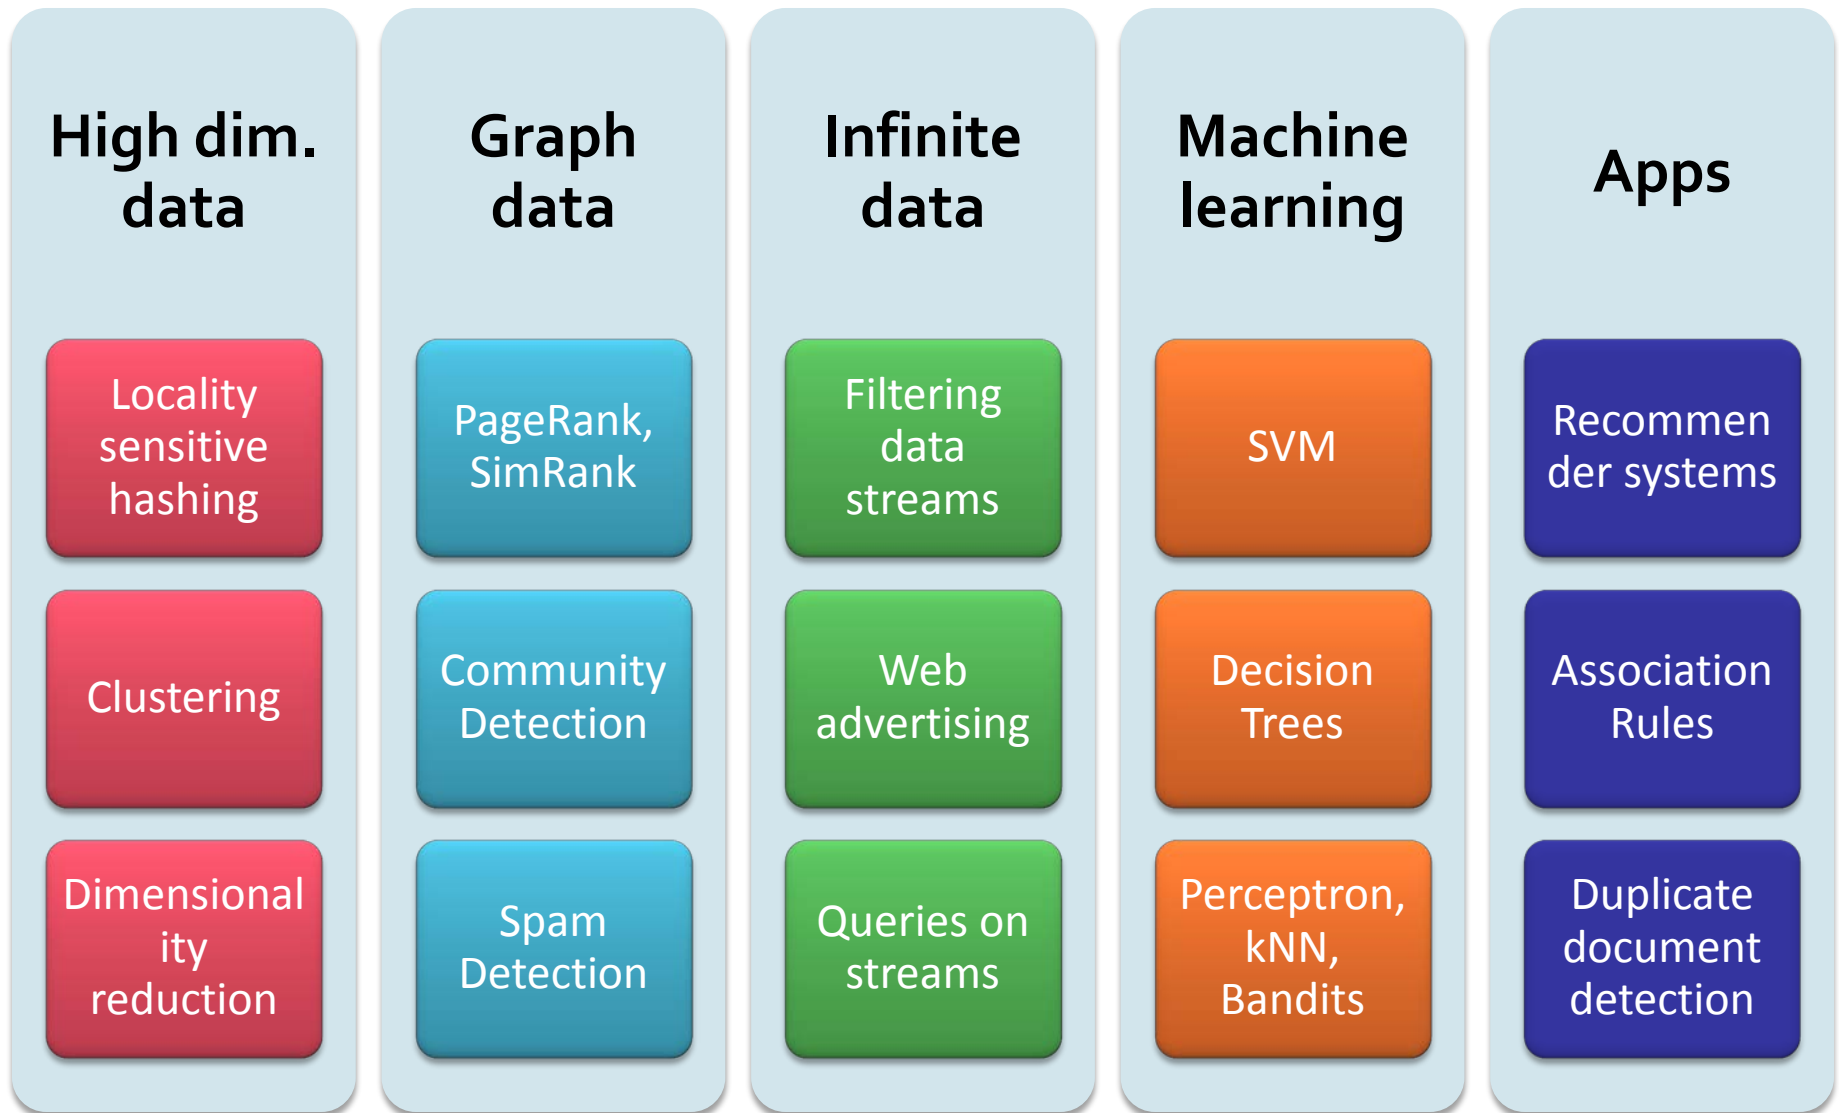
How It All Fits Together

- **Based on different types of data:**
 - Data is **high dimensional**
 - Data is a **graph**
 - Data is **never-ending**
 - Data is **labeled**
- **Based on different models of computation:**
 - **Single machine in-memory**
 - **MapReduce**
 - **Streams**
 - **Batch (offline) vs. Active (online) algorithms**

How It All Fits Together

- **Based on different applications:**
 - Recommender systems
 - Market basket analysis
 - Link analysis, spam detection
 - Duplicate detection and similarity search
 - Web advertising
- **Based on different “tools”:**
 - Linear algebra: SVD, Matrix factorization
 - Optimization: Stochastic gradient descent
 - Dynamic programming: Frequent itemsets
 - Hashing: LSH, Bloom filters

How It All Fits Together



How it all fits together?

Data is High-dimensional:

- Locality Sensitive Hashing
- Dimensionality reduction
- Clustering

Data is a graph:

- Link Analysis: PageRank, TrustRank, Hubs & Authorities

Data is Labeled (Machine Learning):

- kNN, Perceptron, SVM, Decision Trees

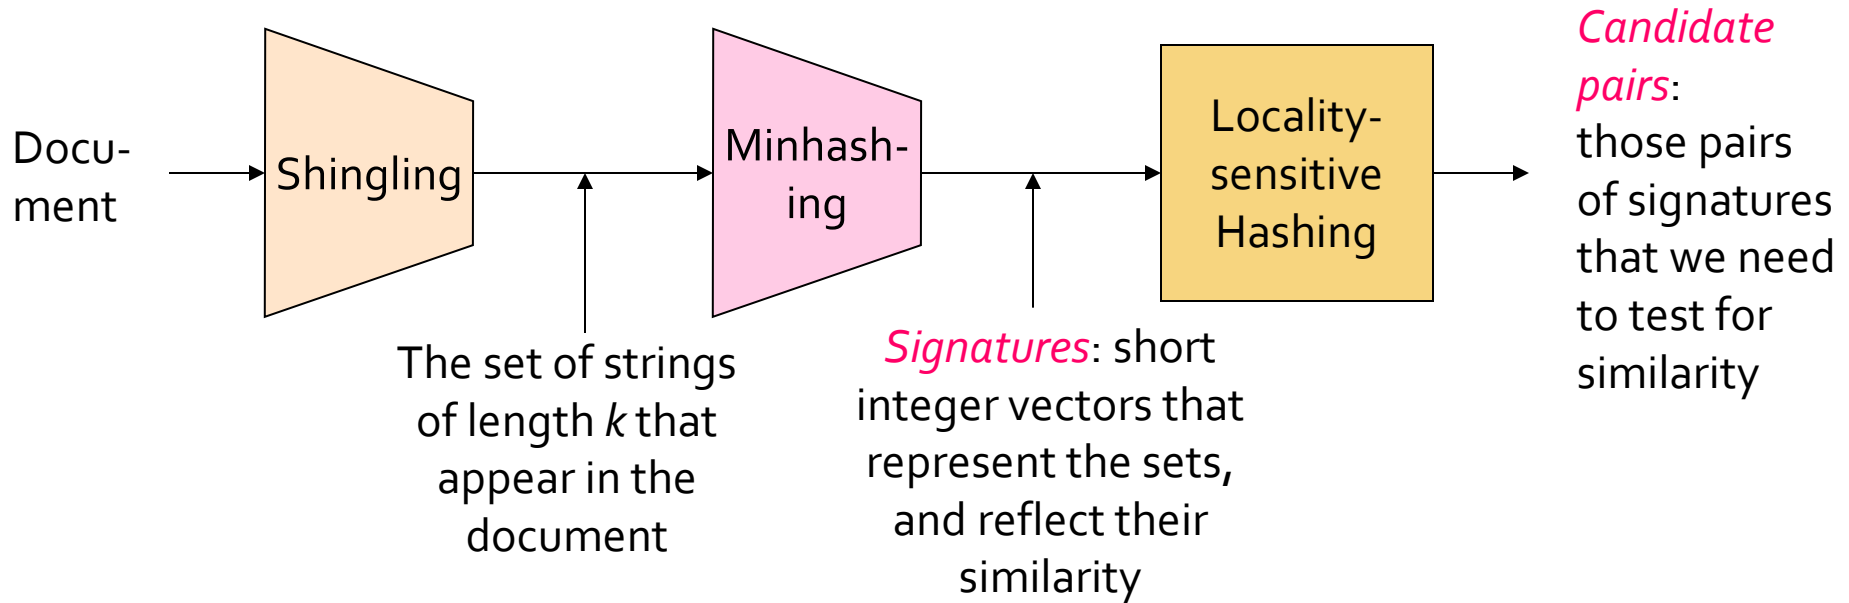
Data is infinite:

- Mining data streams
- Advertising on the Web

Applications:

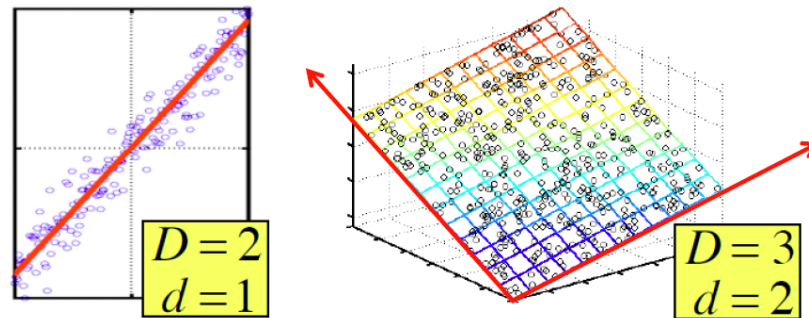
- Association Rules
- Recommender systems

(1) Finding “similar” sets

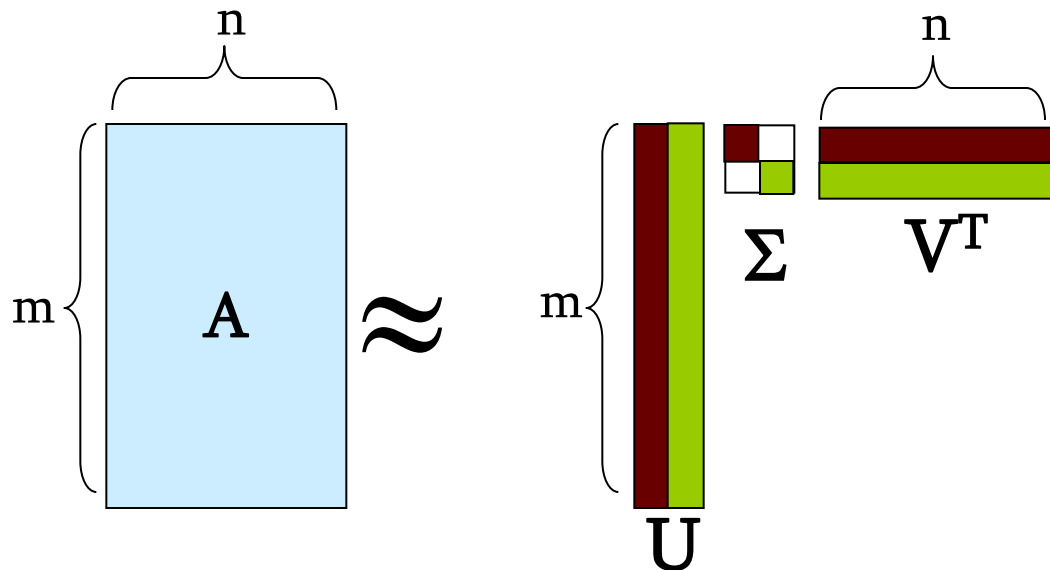


1. **Shingling**: Convert docs to sets
2. **Minhashing**: Convert large sets to short signatures, while preserving similarity
3. **Locality-sensitive hashing**: Focus on pairs of signatures likely to be of similar documents

(2) Dimensionality Reduction



$$\mathbf{A} \approx \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T = \sum_i \sigma_i \mathbf{u}_i \circ \mathbf{v}_i$$

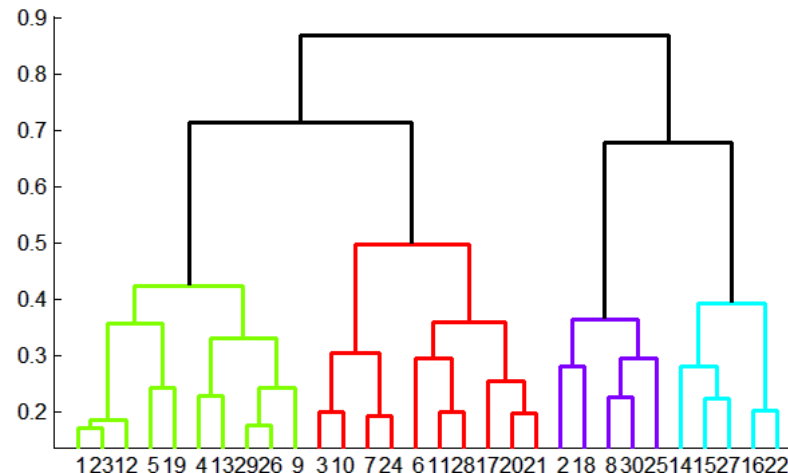


(3) Clustering

■ Hierarchical:

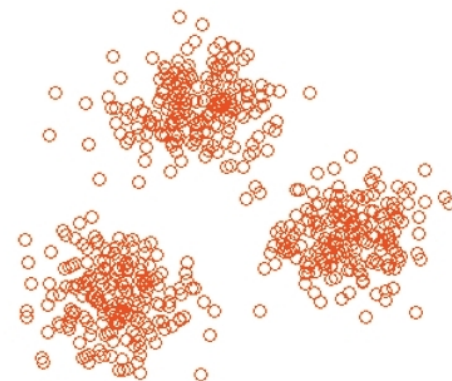
■ **Agglomerative** (bottom up):

- Initially, each point is a cluster
- Repeatedly combine the two “nearest” clusters into one
- Represent a cluster by its **centroid** or **clustroid**



■ **Point Assignment: k-means, BFR**

- Maintain a set of clusters
- Points belong to “nearest” cluster



High-dim data methods: Comparison

■ LSH:

- Find **somewhat** similar pairs of items while avoiding $O(N^2)$ comparisons

■ Clustering:

- Assign points into a **pre-specified** number of **clusters**
 - Each point belongs to a single cluster
 - Summarize the cluster by a centroid

■ SVD (dimensionality reduction):

- Want to explore/exploit **correlations** in the data
- Some dimensions may be irrelevant
- Useful for visualization, removing noise from the data, detecting anomalies

When to use which method?

- **Find all similar pairs of items: LSH**
 - Have to know the threshold ahead of time
 - Allow for some error
- **Identify clusters (structure in data): k-means**
 - k is usually relatively small ($10 \sim 1000$)
 - Useful for identifying ‘types’ or ‘classes’ of datapoints
- **Build low-dimensional representation of data: SVD**
 - More robust (noise-free) similarity computation
 - Data compression (memory saving, speed-up)

How it all fits together?

Data is high-dimensional:

- Locality Sensitive Hashing
- Dimensionality reduction
- Clustering

The data is a graph:

- Link Analysis: PageRank, TrustRank, Hubs & Authorities

Data is labeled (Machine Learning):

- kNN, Perceptron, SVM, Decision Trees

Data is infinite:

- Mining data streams
- Advertising on the Web

Applications:

- Association Rules
- Recommender systems

Link Analysis: PageRank

- Rank nodes using the network link structure

- PageRank:

- Link voting:

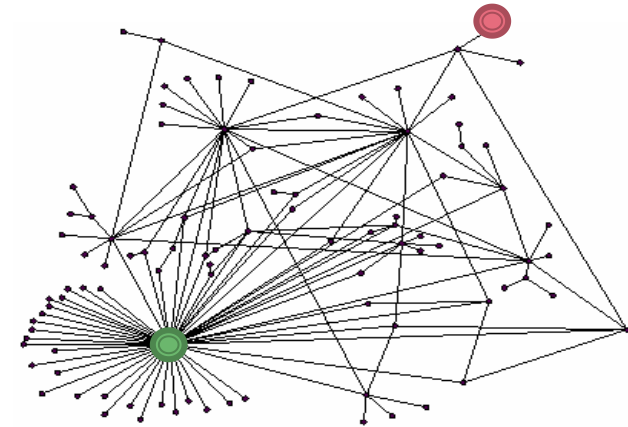
- Page of importance x has n out-links, each gets x/n votes
 - Page R 's importance is the sum of the votes on its in-links

- Complications: Spider traps, Dead-ends

- Solution: At each step, random surfer has 2 options

- With probability β , follow a link at random
 - With prob. $1-\beta$, jump to some page **uniformly** at random

- Power method to compute PageRank



PPR, SimRank, HITS

- **Personalized (topic specific) PageRank**

- Random walker teleports to a preselected set of nodes

- **Random Walk with Restarts**

- Random walker always jumps back to the starting node

- **SimRank**

- **Measure similarity between items**

- k -partite graph with k types of nodes

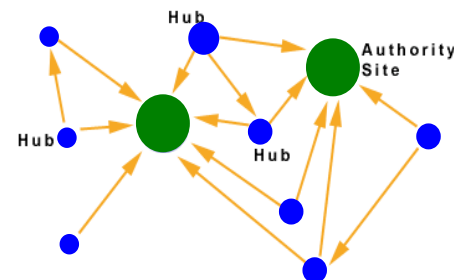
- Perform a random-walk with restarts from node N

- Resulting prob. distrib. is similarity of other nodes to N

- **Hubs & Authorities**

- **Experts vs. Content providers**

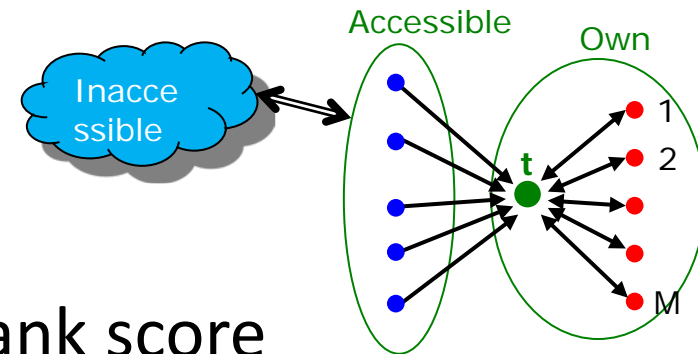
- Principle of repeated improvement



WebSpam and PageRank

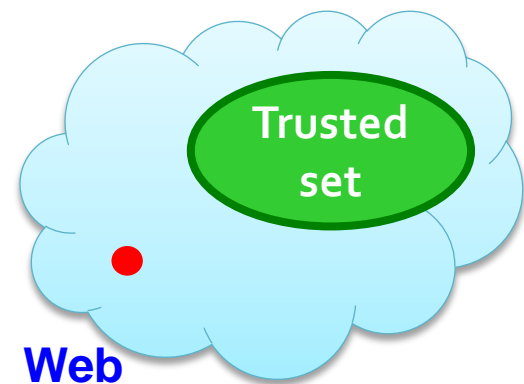
■ Web spam farming

- Architecture of a spam farm
- Effect of spam farms on PageRank score



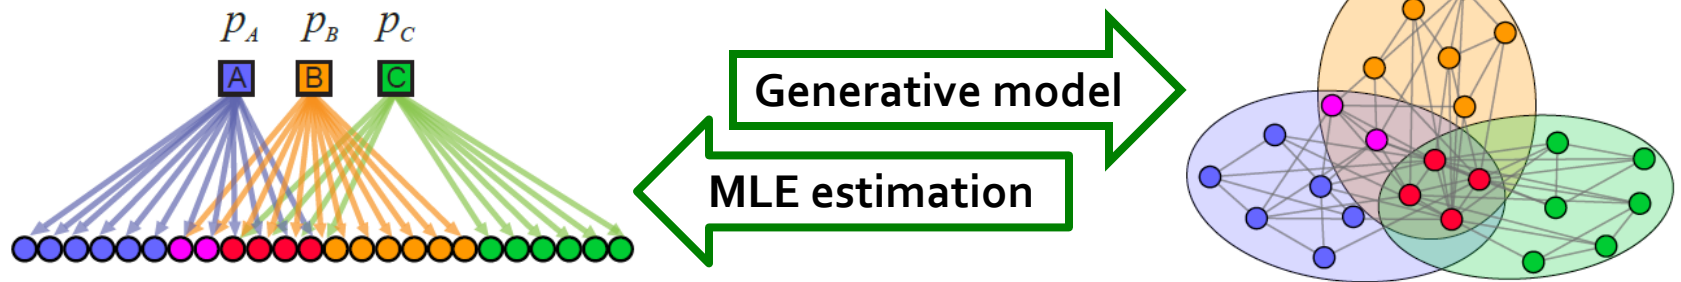
■ TrustRank

- Topic specific PageRank with a teleport set of “trusted” pages
- Spam Mass of a page

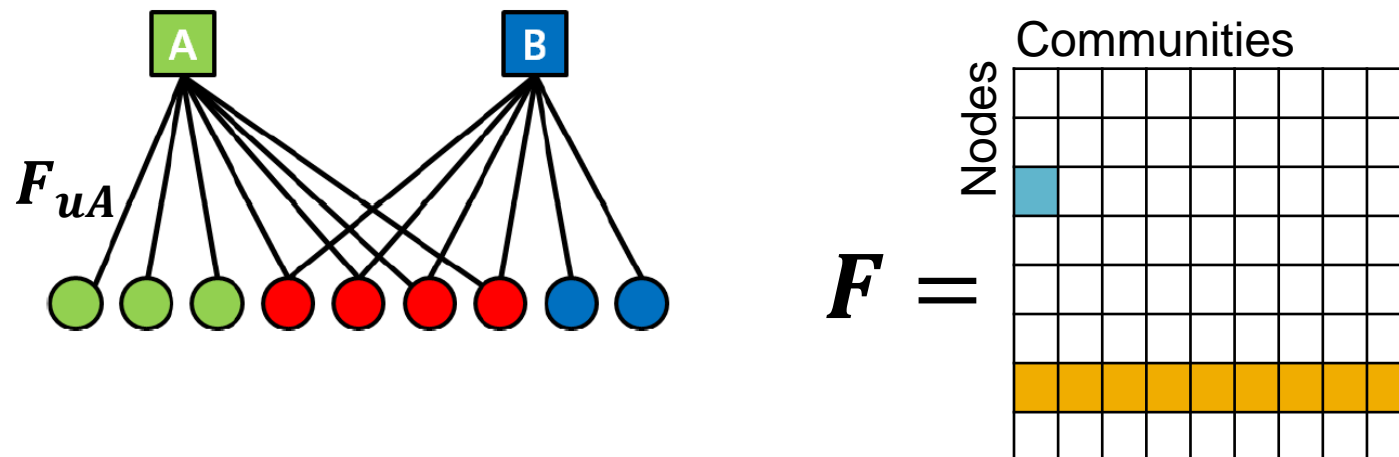


Analysis of Large Graphs

■ AGM (Affiliation Graph Model)



■ BigCLAM (CLuster Affiliation Model)



How it all fits together?

Data is high-dimensional:

- Locality Sensitive Hashing
- Dimensionality reduction
- Clustering

The data is a graph:

- Link Analysis: PageRank, TrustRank, Hubs & Authorities

Data is labeled (Machine Learning):

- kNN, Perceptron, SVM, Decision Trees

Data is infinite:

- Mining data streams
- Advertising on the Web

Applications:

- Association Rules
- Recommender systems

Support Vector Machines

- **Prediction = $\text{sign}(w \cdot x + b)$**

- Model parameters w, b

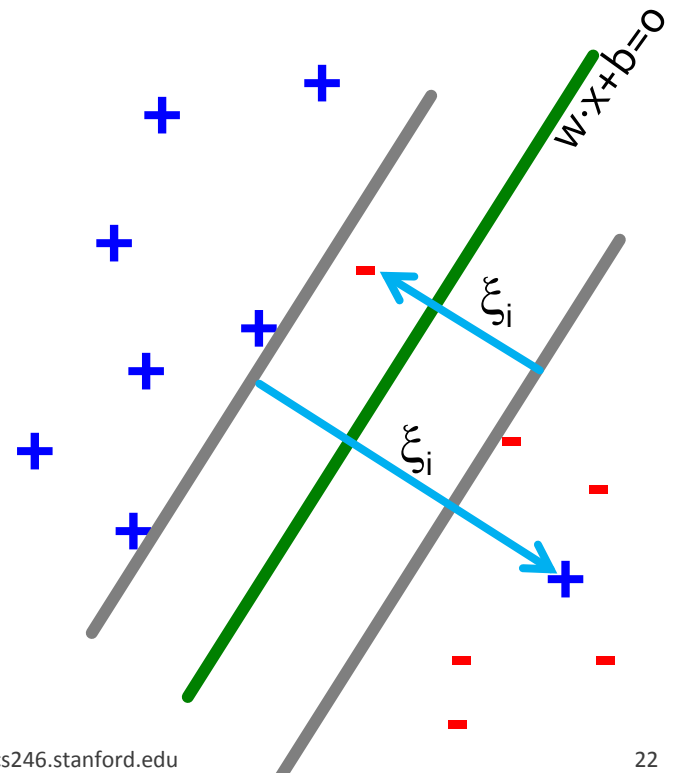
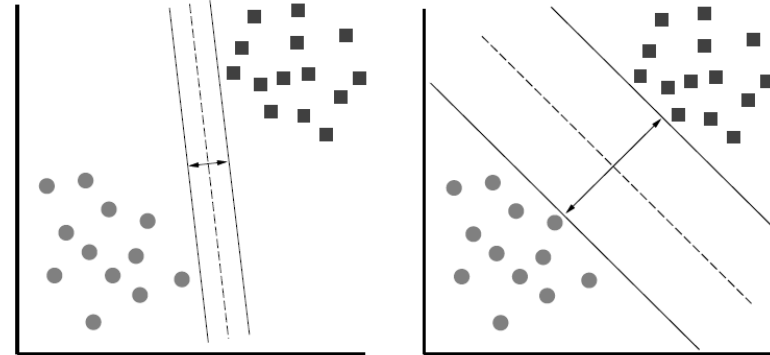
- **Margin:** $\gamma = \frac{\|w\|}{w \cdot w} = \frac{1}{\|w\|}$

- **SVM optimization problem:**

$$\min_{w, b, \xi_i \geq 0} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^n \xi_i$$

$$s.t. \forall i, y_i (w \cdot x_i + b) \geq 1 - \xi_i$$

- Find w, b using **Stochastic gradient descent**



Decision Trees

■ Building decision trees using MapReduce

■ How to predict?

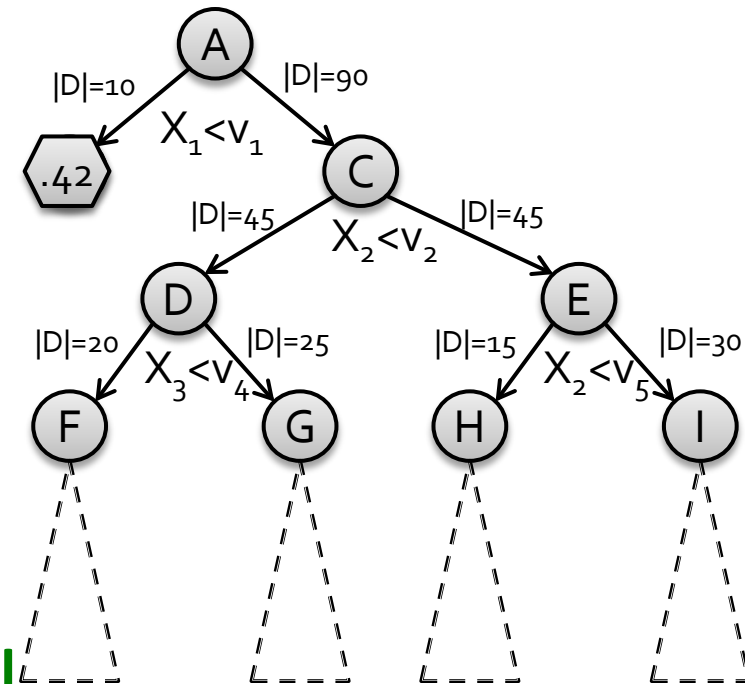
- **Predictor:** avg. y_i of the examples in the leaf

■ When to stop?

- # of examples in the leaf is small

■ How to build?

- One MapReduce job per level
 - Need to compute split quality for each attribute and each split value for each current leaf



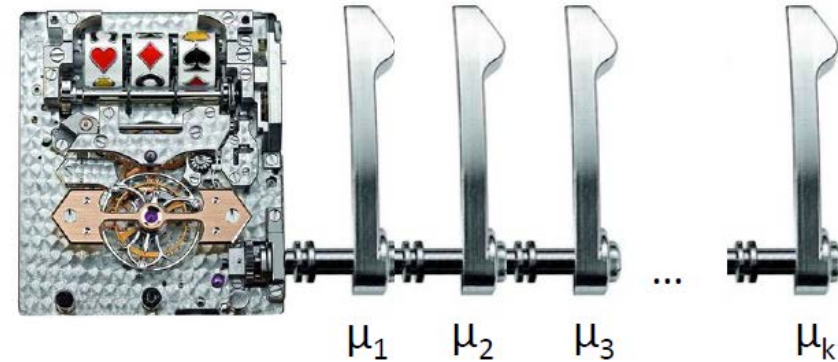
Algorithm 1 FindBestSplit

Require: Node n , Data $D \subseteq D^*$

- 1: $(n \rightarrow \text{split}, D_L, D_R) = \text{FindBestSplit}(D)$
- 2: if StoppingCriteria(D_L) then
- 3: $n \rightarrow \text{left_prediction} = \text{FindPrediction}(D_L)$
- 4: else
- 5: $\text{FindBestSplit}(n \rightarrow \text{left}, D_L)$
- 6: if StoppingCriteria(D_R) then
- 7: $n \rightarrow \text{right_prediction} = \text{FindPrediction}(D_R)$
- 8: else
- 9: $\text{FindBestSplit}(n \rightarrow \text{right}, D_R)$

Learning Through Experimentation

- **Learning through experimentation**
 - Exploration-Exploitation tradeoff
 - Regret
- **Multiarmed Bandits**
 - Epsilon-Greedy
 - UCB1 algorithm
- **Submodular function optimization**
 - Coverage
 - Greedy and Lazy-Greedy algorithms
 - Multiplicative Weights algorithm



When to use which method?

- **SVM: Classification**
 - **Millions of sparse numerical features** (e.g., documents)
 - Simple (linear) decision boundary
 - Somewhat hard to interpret model
- **k-NN: Classification or regression**
 - (Many) numerical features
 - Many design decisions – distance metric, k , weighting, ...
there is no simple way to set them!
- **Decision Trees: Classification or Regression**
 - Relatively few dense features (handles categorical features)
 - Complicated decision boundary: **Overfitting!**
 - **Easy to explain/interpret the classification**
 - **Bagged Decision Trees** – very, very hard to beat!
- **Bandits: Learning through experimentation**
 - Exploration-Exploitation tradeoff

What if “ML alg. doesn’t work”?

- **Over- vs. under-fitting**

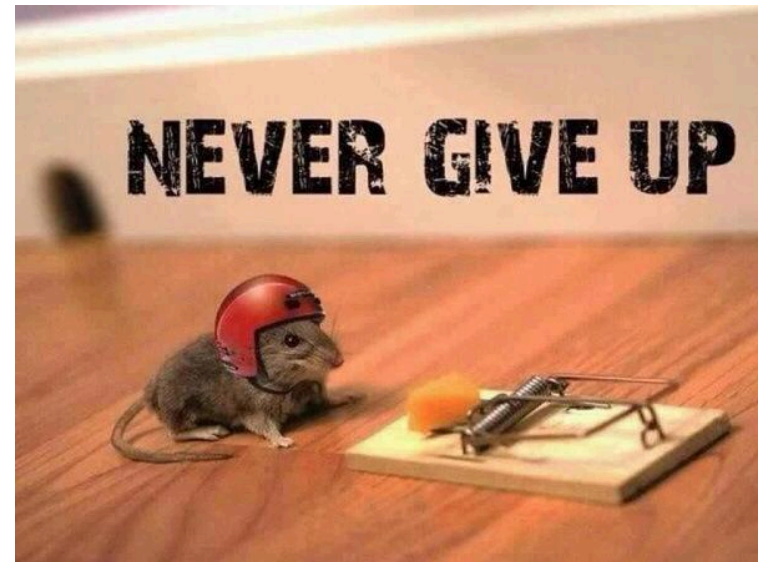
- Compare error on the train/test set
- Plot error vs. (regularization) parameter

- **Debugging:**

- Compare performance to a simple baseline
- Build synthetic datasets for which you know your method should work

- **Think about:**

- The prediction problem
- Error metrics
- Model assumptions
- Properties of the data



If the ML algorithm doesn't work

- **Get more training data**
 - Sometimes more data doesn't help but often it does
- **Try a smaller set a features**
 - Carefully select small subset
 - You can do this by hand, or use SVD
- **Try getting additional features**
 - **LOOK** at the data
 - Can be very time consuming
- **Adding polynomial features**
 - Include \mathbf{x} and \mathbf{x}^2 as features
- **Building your own, new, better features**
 - Based on your knowledge of the problem
- **Try decreasing or increasing regularization parameter**
 - Change how important the regularization term is

How it all fits together?

Data is high-dimensional:

- Locality Sensitive Hashing
- Dimensionality reduction
- Clustering

The data is a graph:

- Link Analysis: PageRank, TrustRank, Hubs & Authorities

Data is labeled (Machine Learning):

- kNN, Perceptron, SVM, Decision Trees

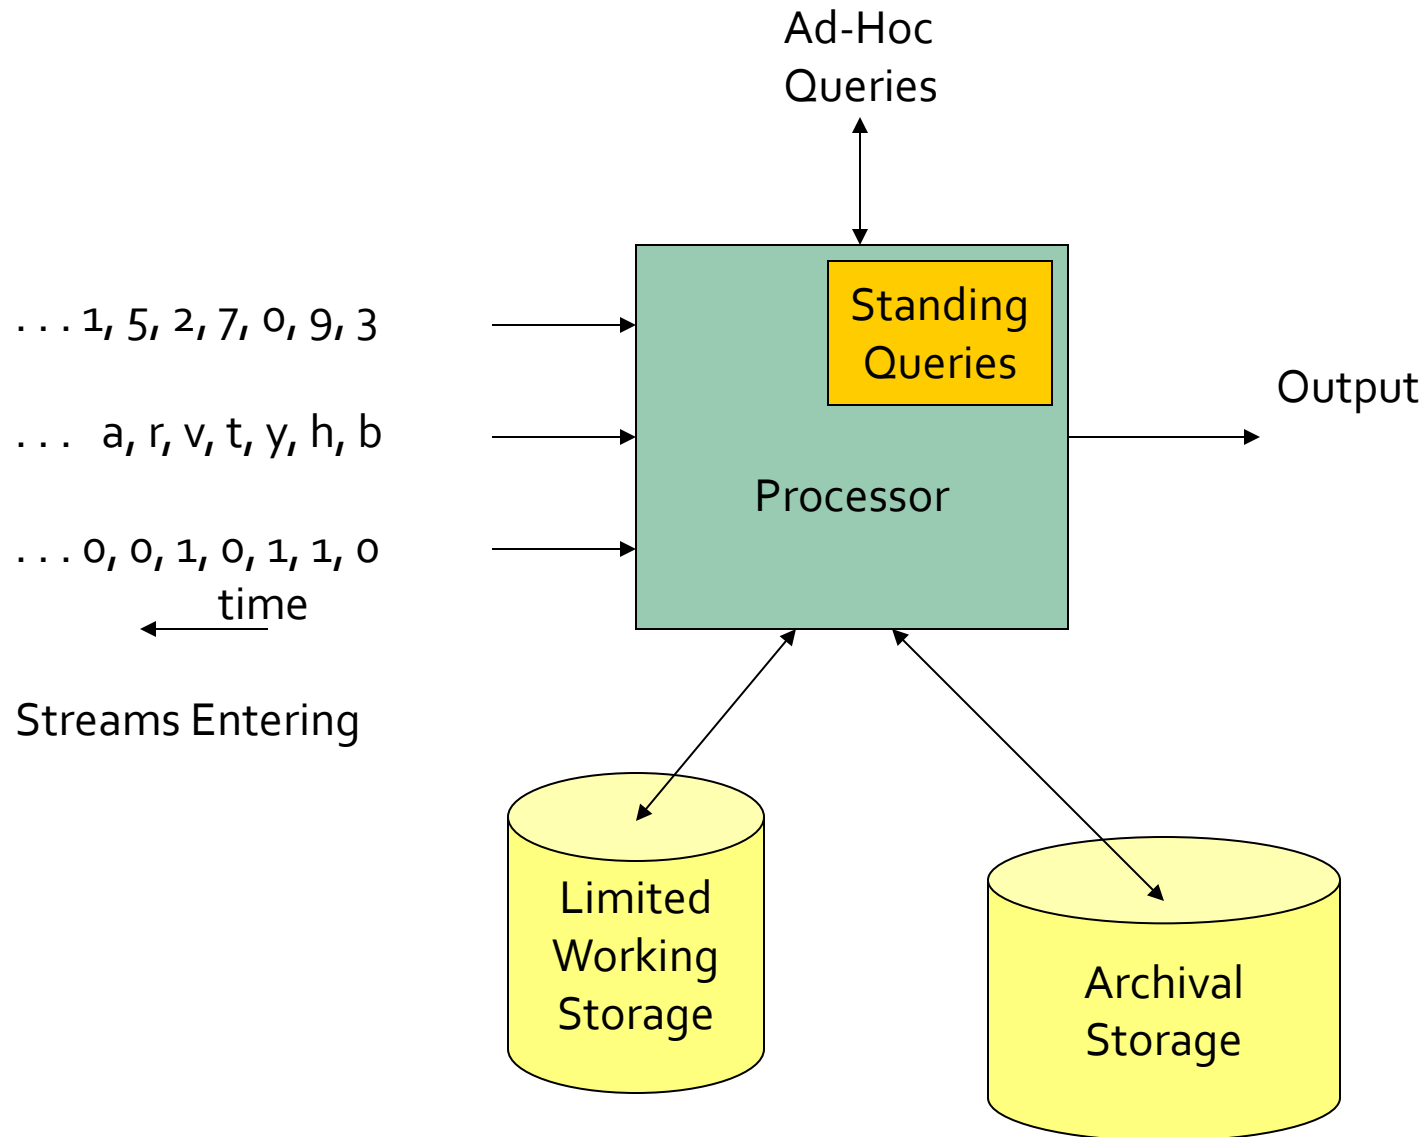
Data is infinite:

- Mining data streams
- Advertising on the Web

Applications:

- Association Rules
- Recommender systems

Mining Data Streams



Problems on data streams

- **Sampling data from a stream:**

- Each element is included with prob. k/N

- **Queries over sliding windows:**

How many **1**s are in last k bits?

1001010110001011010101010101011010101010101110101010111010101011101010100010110010

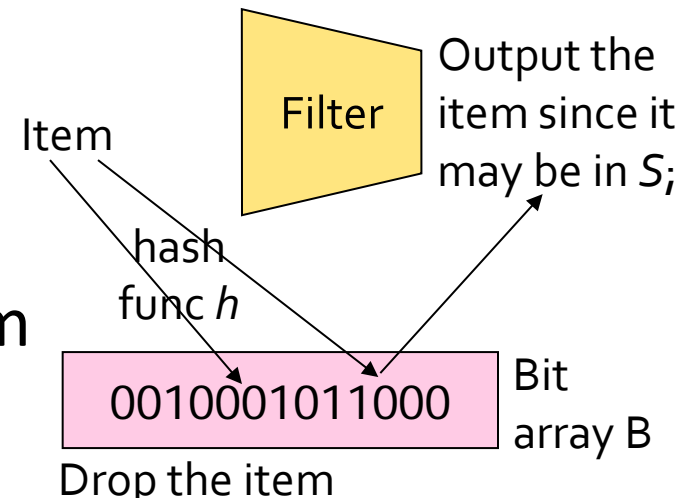
- **Filtering a stream: Bloom filters**

- Filter elements with property x

- **Counting distinct elements:**

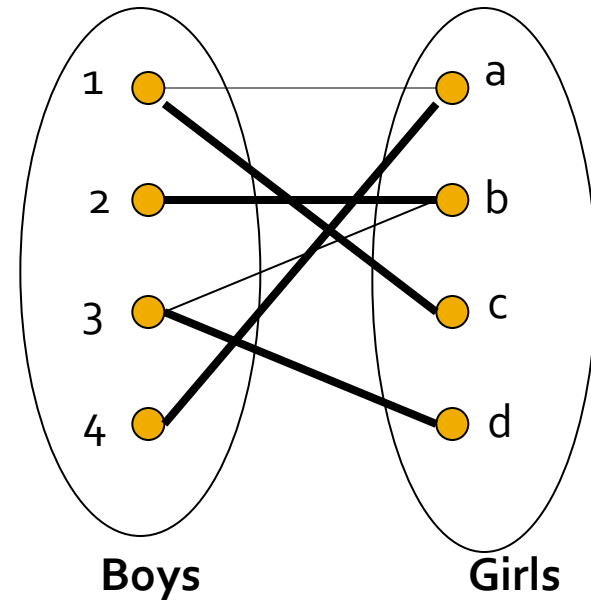
- Number of distinct elements in the last k elements of the stream

- **Estimating moments**



Online algorithms & Advertising

- You get to see one input piece at a time, and need to make irrevocable decisions
- **Competitive ratio** = $\min_{\text{all inputs}} (|M_{\text{my_alg}}| / |M_{\text{opt}}|)$
- **Adwords problem:**
 - Query arrives to a search engine
 - Several advertisers bid on the query
 - Pick a subset of advertisers whose ads are shown
- **Greedy online matching: competitive ratio $\geq 1/2$**



How it all fits together?

Data is high-dimensional:

- Locality Sensitive Hashing
- Dimensionality reduction
- Clustering

The data is a graph:

- Link Analysis: PageRank, TrustRank, Hubs & Authorities

Data is labeled (Machine Learning):

- kNN, Perceptron, SVM, Decision Trees

Data is infinite:

- Mining data streams
- Advertising on the Web

Applications:

- Association Rules
- Recommender systems

Association Rule Discovery

Market-basket model:

- **Goal:** To identify items that are bought together by sufficiently many customers
- **Approach:** Process the sales data collected with barcode scanners to find dependencies among items

Discovering frequent items: A-priori, PCY

<i>TID</i>	<i>Items</i>
1	Bread, Coke, Milk
2	Beer, Bread
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Coke, Diaper, Milk

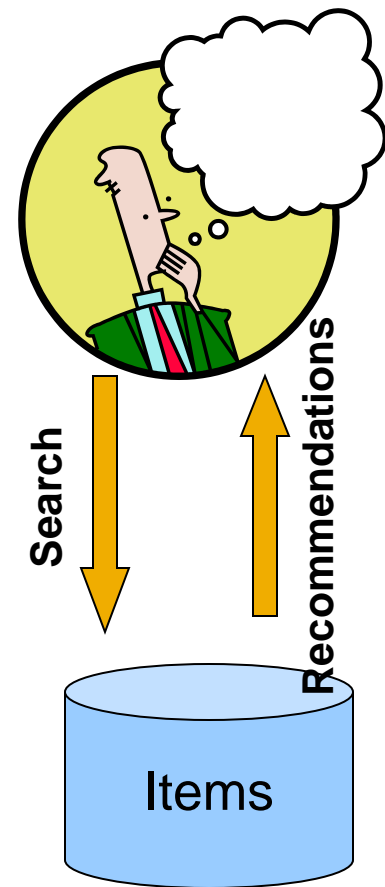
Rules Discovered:

$\{\text{Milk}\} \rightarrow \{\text{Coke}\}$

$\{\text{Diaper, Milk}\} \rightarrow \{\text{Beer}\}$

Recommender Systems

- **User-user collaborative filtering**
 - Consider user c
 - Find set D of other users whose ratings are “similar” to c ’s ratings
 - Estimate user’s ratings based on the ratings of users in D
- **Item-item collaborative filtering**
 - Estimate rating for item based on ratings for similar items
- **Profile based**



Latent Factor Models: Netflix

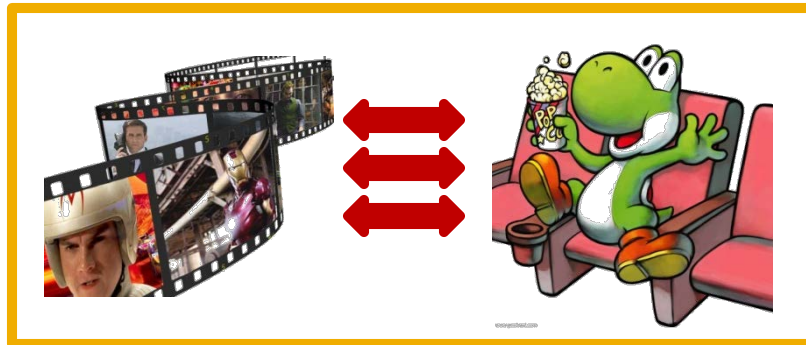
user bias



movie bias



user-movie interaction



Baseline predictor

- Separates users and movies
- Benefits from insights into user's behavior

User-Movie interaction

- Characterizes the matching between users and movies
- Attracts most research in the field

$$\min_{Q, P} \sum_{(u, i) \in R} \left(r_{ui} - (\mu + b_u + b_i + q_i^T p_u) \right)^2$$
$$+ \lambda \left(\sum_i \|q_i\|^2 + \sum_u \|p_u\|^2 + \sum_u \|b_u\|^2 + \sum_i \|b_i\|^2 \right)$$

When to use which method?

- **Lots of rating data: CF**
 - Easy to tweak, easy to add lots of features/signals
 - Use optimization to learn weights on how to combine features
- **Lots² of rating data: CF + Latent factors**
 - Many ratings per user, many ratings per item
 - Depending on the amount of data make the model more/less complex (more/less parameters)
- **Cold start, little data: Profile based**
 - Need to have good user/item features and similarity metric

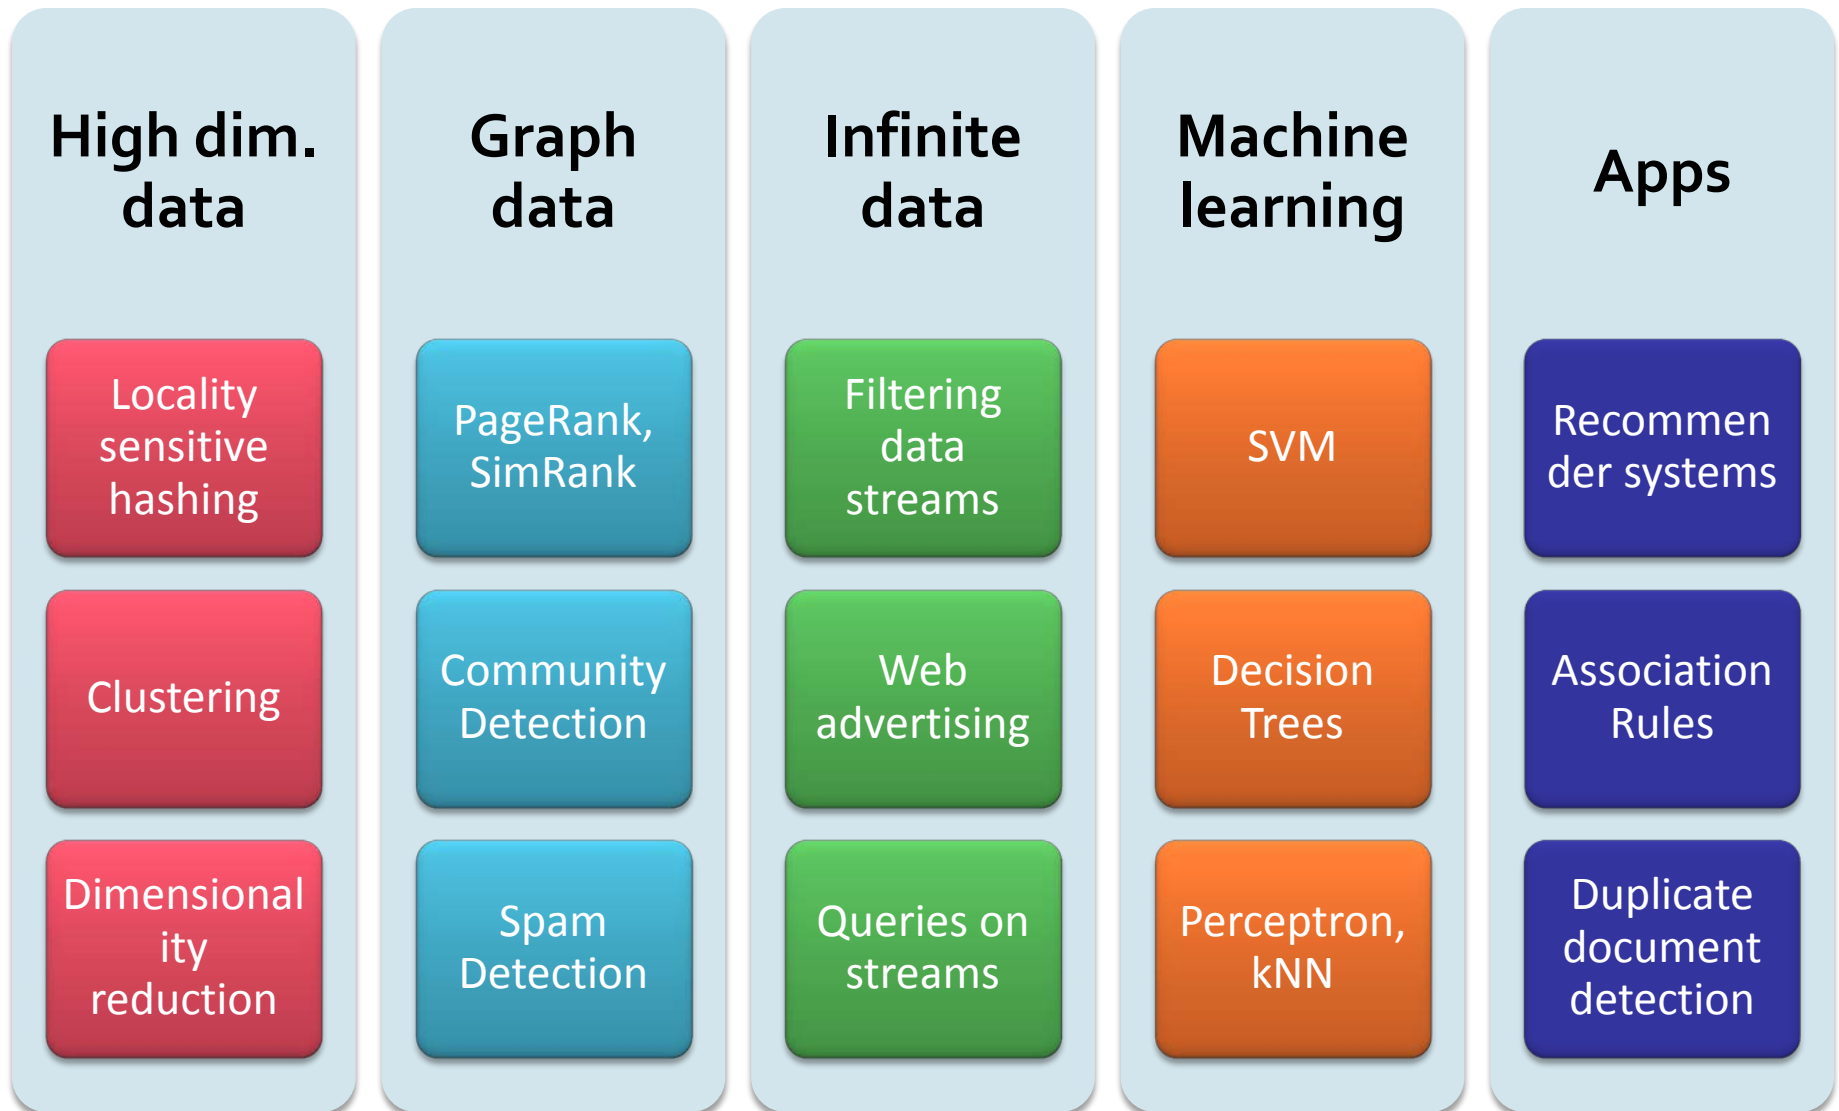
In closing...

What we've learned this quarter

- MapReduce
- Association Rules
- Apriori algorithm
- Finding Similar Items
- Locality Sensitive Hashing
- Random Hyperplanes
- Dimensionality Reduction
- Singular Value Decomposition
- CUR method
- Clustering
- Recommender systems
- Collaborative filtering
- PageRank and TrustRank
- Hubs & Authorities
- k-Nearest Neighbors
- Perceptron
- Support Vector Machines
- Stochastic Gradient Descent
- Decision Trees
- Mining data streams
- Bloom Filters
- Flajolet-Martin
- Advertising on the Web



Map of Superpowers



Applying Your Superpowers



THE BIG PICTURE

- How to analyze large datasets to discover **models** and **patterns** that are:
 - **Valid:** Hold on new data with some certainty
 - **Novel:** Non-obvious to the system
 - **Useful:** Should be possible to act on the item
 - **Understandable:** Humans should be able to interpret the pattern

What next? Seminars

■ Seminars:

- InfoSeminar: <http://i.stanford.edu/infoseminar>
- RAIN Seminar: <http://rain.stanford.edu>

■ Conferences:

- **KDD**: ACM Conf. on Knowledge Discovery & Data Mining
- **WSDM**: ACM Conf. on Web Search and Data Mining
- **ICDM**: IEEE International Conf. on Data Mining
- **WWW**: World Wide Web Conference
- **ICML**: International Conf. on Machine Learning
- **VLDB**: Very Large Data Bases

CS341: Project in Data Mining

- **Data mining research project on real data**
 - Groups of 3 students
 - **We provide interesting data, computing resources (Amazon EC2) and mentoring**
 - **You provide project ideas**
 - There are (practically) no lectures, only individual group mentoring

Information session:
Today 6pm in Gates 415
(there will be pizza)

What Next? Courses

- **Other relevant courses**
 - **CS224W**: Social and Information Network Analysis
 - **CS276**: Information Retrieval and Web Search
 - **CS229**: Machine Learning
 - **CS245**: Database System Principles
 - **CS347**: Distributed Databases
 - **CS448g**: Interactive Data Analysis

What Next? Final Exam



In Closing

- **You Have Done a Lot!!!**
- **And (hopefully) learned a lot!!!**
 - Answered questions and proved many interesting results
 - Implemented a number of methods
 - **And did excellently on the final!**

**Thank You for the
Hard Work!!!**